

Distributed databases building blocks

OREN EINI

WIZARD

HIBERNATING RHINOS

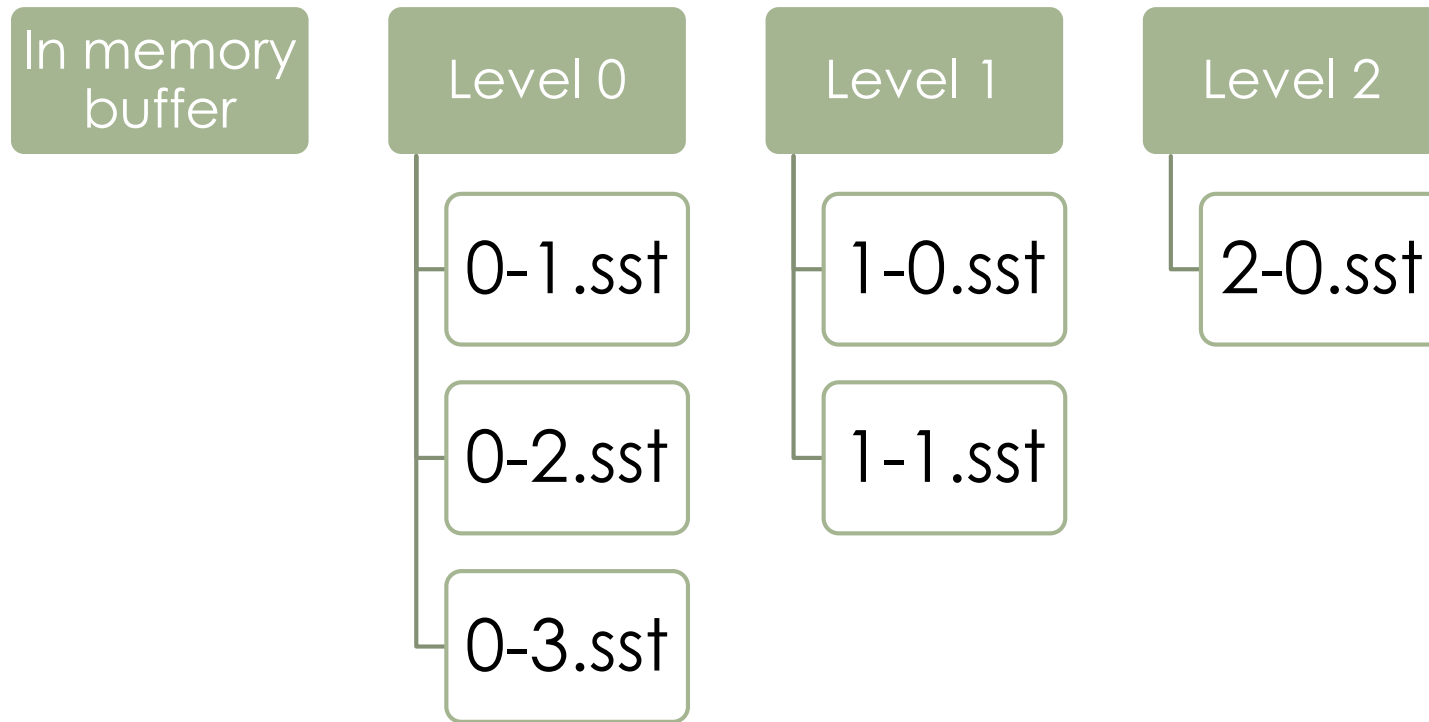
I write databases for a living

- ▶ Rhino DHT
- ▶ RavenDB
 - ▶ Voron
- ▶ This is hard 😊
 - ▶ Also fun

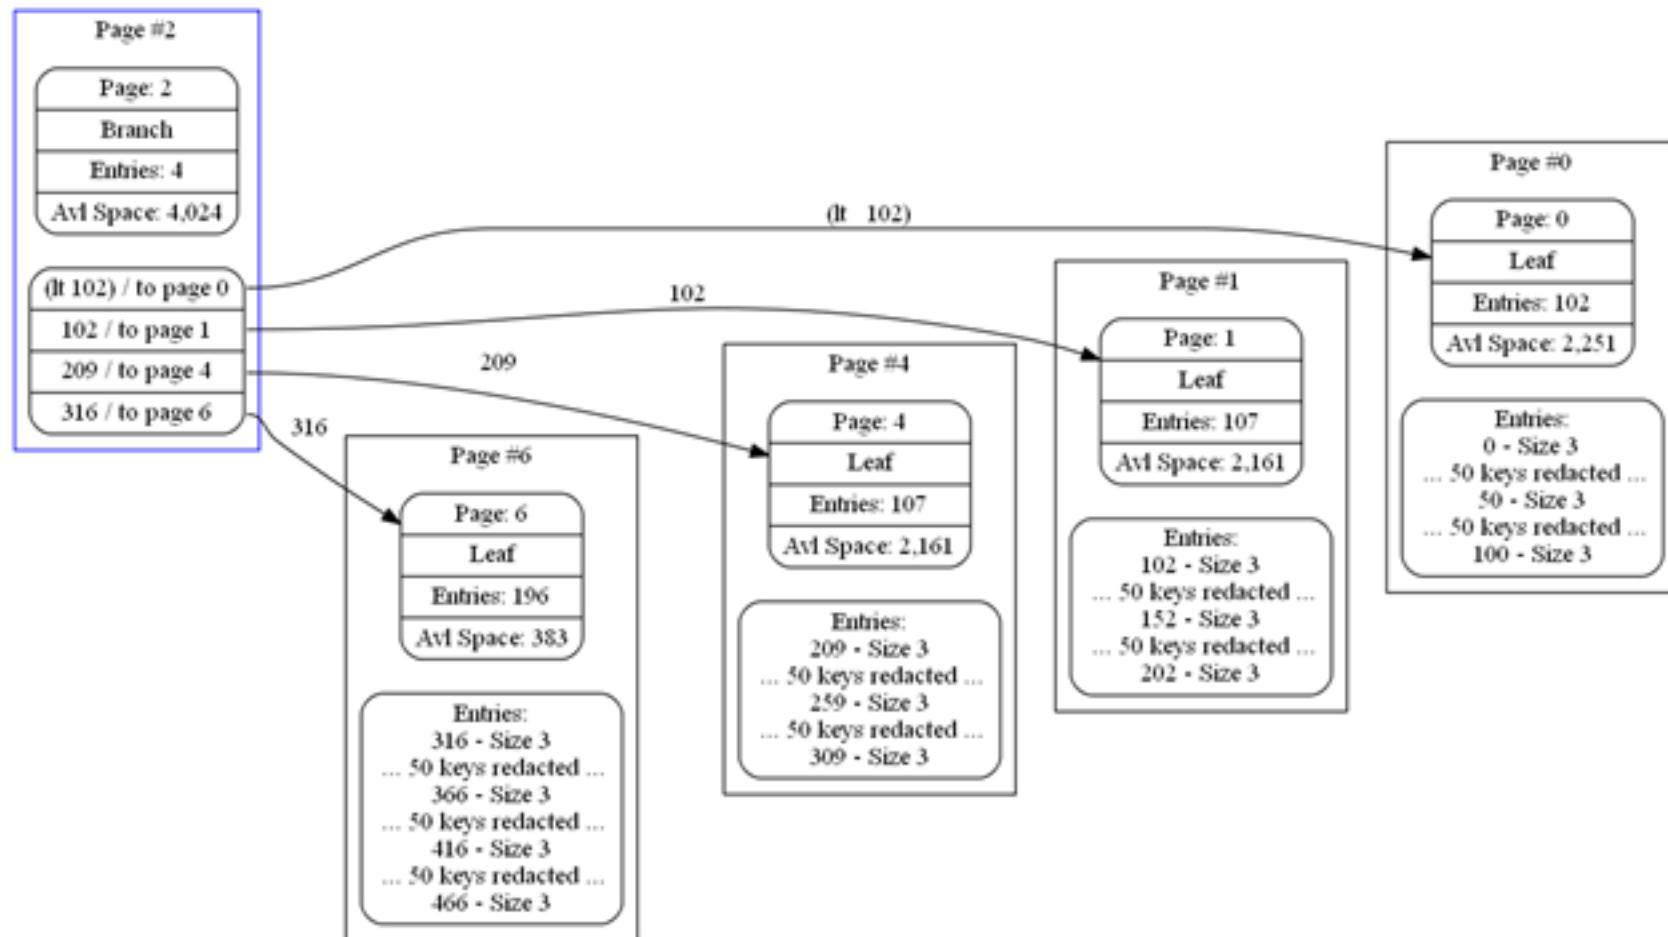
Storing data

- ▶ What do we need?
- ▶ File.Open ?
- ▶ The disk goes round & round, round & round, round & round

Log Structure Merge



B+Trees



Concurrency & Isolations

Tx 1


- ▶ Begin Tx
- ▶ Read X (X = 1)
- ▶ Write X = 2
- ▶ Commit

Tx 2

- ▶ Begin Tx
- ▶ Read X (X = 1)
- ▶ Read X (X = ?)

Tx 3

- ▶ Begin Tx
- ▶ Read X (X = 2)



Implementing concurrency


Locks



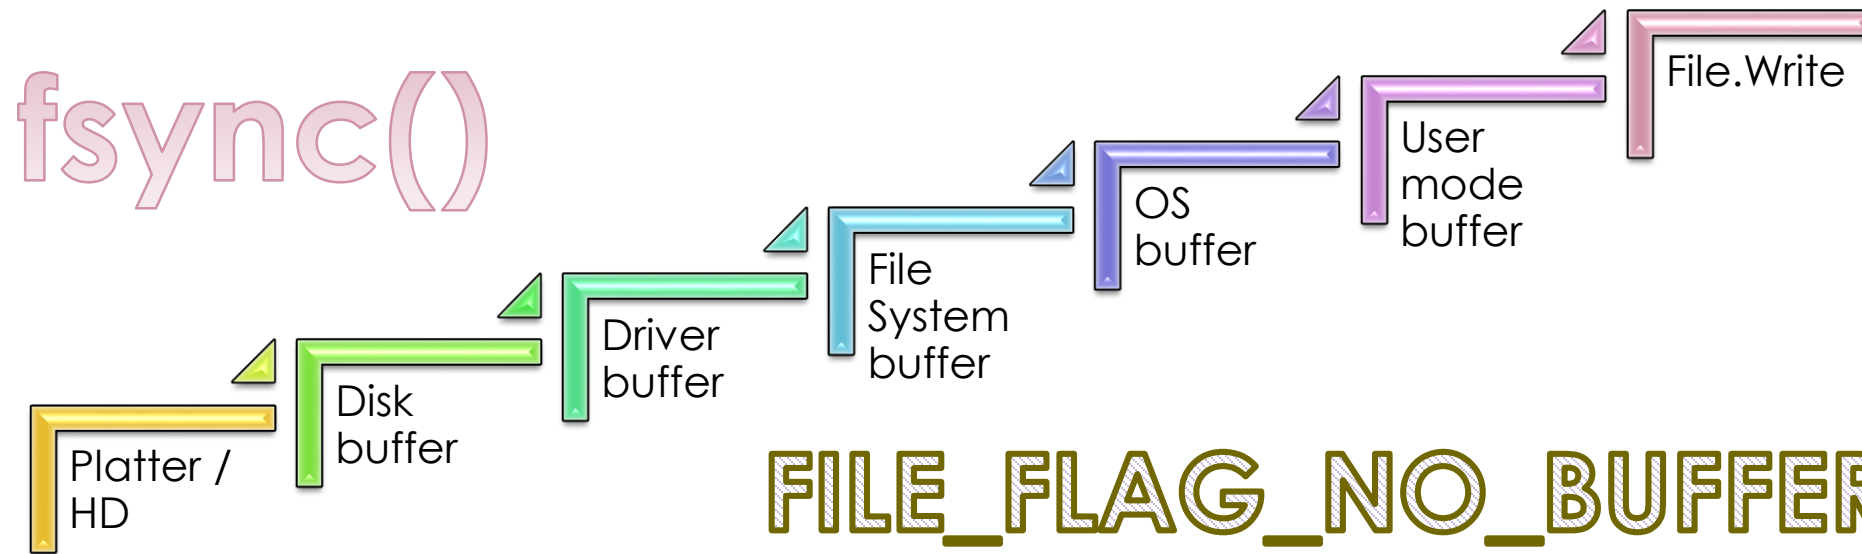
MVCC



I/O IS SLOW



Durability / transactions / fsync




O_SYNC

**FILE_FLAG_NO_BUFFERING |
FILE_FLAG_WRITE_THROUGH**

The curse of the single node...



What about the distributed part?



Before that...

- ▶ What is the distribution model?
 - ▶ Consensus?
 - ▶ Collaborative?
 - ▶ Repair?
- ▶ Partition model

Consensus

Paxos ???



Raft




Consensus – log of operations

- ▶ Set $x = 1$
- ▶ Set $y = 2$
- ▶ Set $n = 1$
- ▶ Set $x = y + n$



Problems?



Collaborations (master / master)



- ▶ Allow writes on any node
- ▶ Conflict-free replicated data types ?
- ▶ Merges

Sharded

- ▶ Some data on some node
- ▶ Multiple consensus groups




Gossip



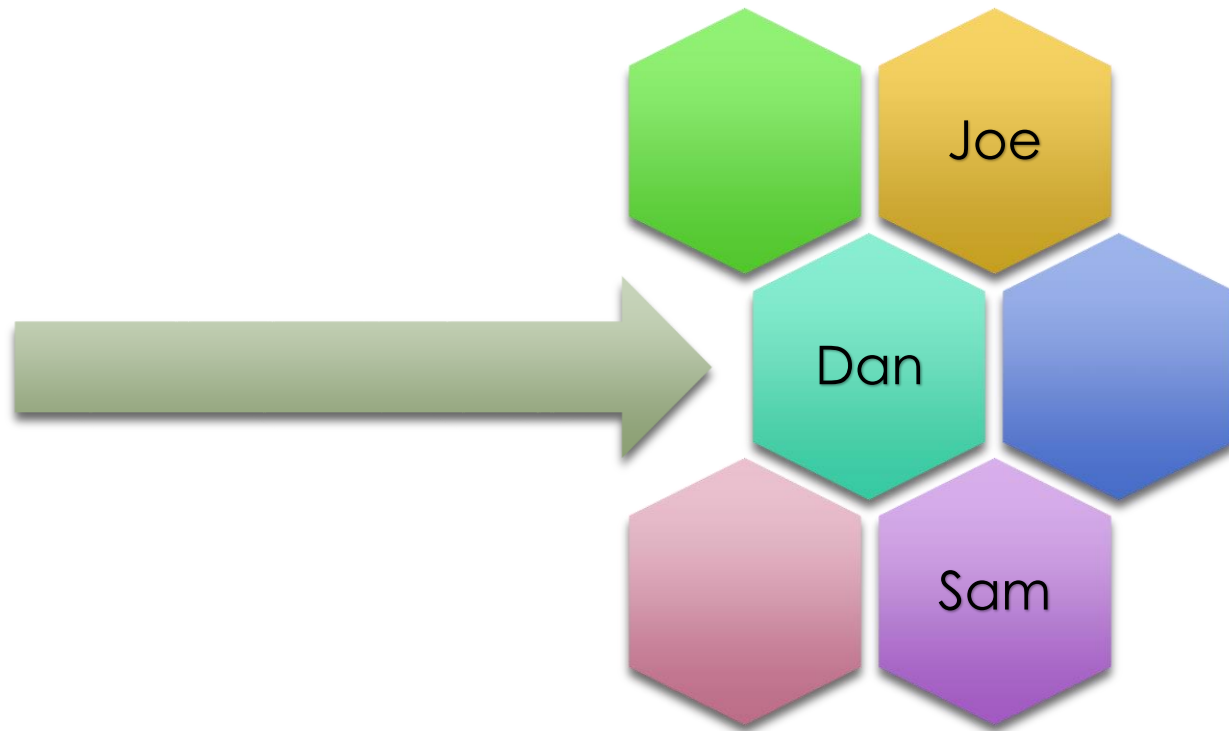
Data model

- ▶ So now you know how to store data (single node, multiple nodes)
- ▶ What do you do with this?

Entire dataset



Portioned data



Imagine a banking system...

- ▶ Core data model:
 - ▶ Accounts
 - ▶ Account operations (incoming, outgoing, fees, etc)
- ▶ New rule changes – gossip
- ▶ Reporting

Questions?

