

# Consistent hashing and distributed hash table

Viet-Trung Tran

[trungtv@soict.hust.edu.vn](mailto:trungtv@soict.hust.edu.vn)

# Outline

# Hashing

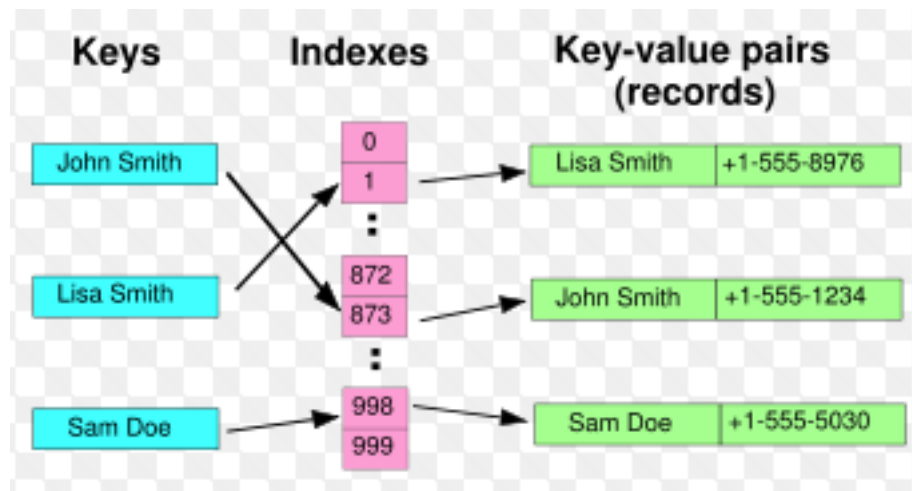
- A data structure for which both searching and insertions are  $O(1)$  in the worst case
- Searching without performing a search ...
  - given an item  $x$ , we need to be able to determine directly from  $x$  the array position where it is to be stored

# Hashing

- The idea of hashing is to distribute the entries of a dataset across an array of buckets.
- Given a key, the algorithm computes an index that suggests where an entry can be found:
  - $\text{index} = f(\text{key}, \text{array\_size})$
- Often this is done in two steps:
  - $\text{hash} = \text{hashfunc}(\text{key})$ .
  - $\text{index} = \text{hash} \% \text{array\_size}$

# A Hash Table (hash map)

- A data structure to implement an associative array.
- A structure that can map keys to values.
- Uses a hash function to compute an index into an array of buckets or slots from which the correct value can be found.



# Hash function

- Crucial for good hash table performance.
- Can be difficult to achieve.
- A basic expectation is that the function would provide a uniform distribution of hash values.
- A non-uniform distribution increases the number of collisions and the cost of resolving them.

# Basic Hashing for Partitioning?

- Consider problem of data partition:
  - Given document  $X$ , choose one of  $k$  servers to use
- Suppose we use modulo hashing
  - Number servers  $1..k$
  - Place  $X$  on server  $i = (X \bmod k)$ 
    - Problem? Data may not be uniformly distributed
  - Place  $X$  on server  $i = \text{hash}(X) \bmod k$ 
    - Problem?
      - What happens if a server fails or joins ( $k \rightarrow k \pm 1$ )?
      - What if different clients have different estimate of  $k$ ?

# What is a DHT?

- Hash Table
  - data structure that maps “keys” to “values”
  - essential building block in software systems
- Distributed Hash Table (DHT)
  - similar, but spread across many hosts
- Interface
  - insert(key, value)
  - lookup(key)




# How do DHTs work?


Every DHT node supports a single operation:

- Given *key* as input; route messages to node holding *key*
  - DHTs are *content-addressable*

# DHT: basic idea




# DHT: basic idea




Neighboring nodes are “connected” at the application-level 11

# DHT: basic idea




Operation: take *key* as input; route messages to node holding *key* **12**

# DHT: basic idea




Operation: take *key* as input; route messages to node holding *key* **13**

# DHT: basic idea




Operation: take *key* as input; route messages to node holding *key* 14

# DHT: basic idea



Operation: take *key* as input; route messages to node holding *key* 15

# DHT: basic idea



Operation: take *key* as input; route messages to node holding *key* 16



# How to design a DHT?

- State Assignment:
  - what “(*key*, *value*) tables” does a node store?
- Network Topology:
  - how does a node select its neighbors?
- Routing Algorithm:
  - which neighbor to pick while routing to a destination?
- Various DHT algorithms make different choices
  - CAN, Chord, Pastry, Tapestry, Plaxton, Viceroy, Kademlia, Skipnet, Symphony, Koorde, Apocrypha, Land, ORDI ...

Taken slides from University of California, berkely and Max planck institute

## **CHORD: A SCALABLE PEER-TO-PEER LOOK-UP PROTOCOL FOR INTERNET APPLICATIONS**

# Outline

- What is Chord?
- Consistent Hashing
- A Simple Key Lookup Algorithm
- Scalable Key Lookup Algorithm
- Node Joins and Stabilization
- Node Failures


# What is Chord?

- In short: a peer-to-peer lookup system
- Given a key (data item), it maps the key onto a node (peer).
- Uses consistent hashing to assign keys to nodes .
- Solves the problem of locating key in a collection of distributed nodes.
- Maintains routing information with frequent node arrivals and departures

# What is Chord? - Addressed Problems

- **Load balance:** chord acts as a distributed hash function, spreading keys evenly over nodes
- **Decentralization:** chord is fully distributed, no node is more important than any other, improves robustness
- **Scalability:** logarithmic growth of lookup costs with the number of nodes in the network, even very large systems are feasible
- **Availability:** chord automatically adjusts its internal tables to ensure that the node responsible for a key can always be found
- **Flexible naming:** chord places no constraints on the structure of the keys it looks up.

# What is Chord? - Typical Application



- Highest layer implements application-specific functions such as file-system meta-data
- This file systems maps operations to lower-level block operations
- Block storage uses Chord to identify the node responsible for storing a block and then talk to the block storage server on that node

# Consistent Hashing

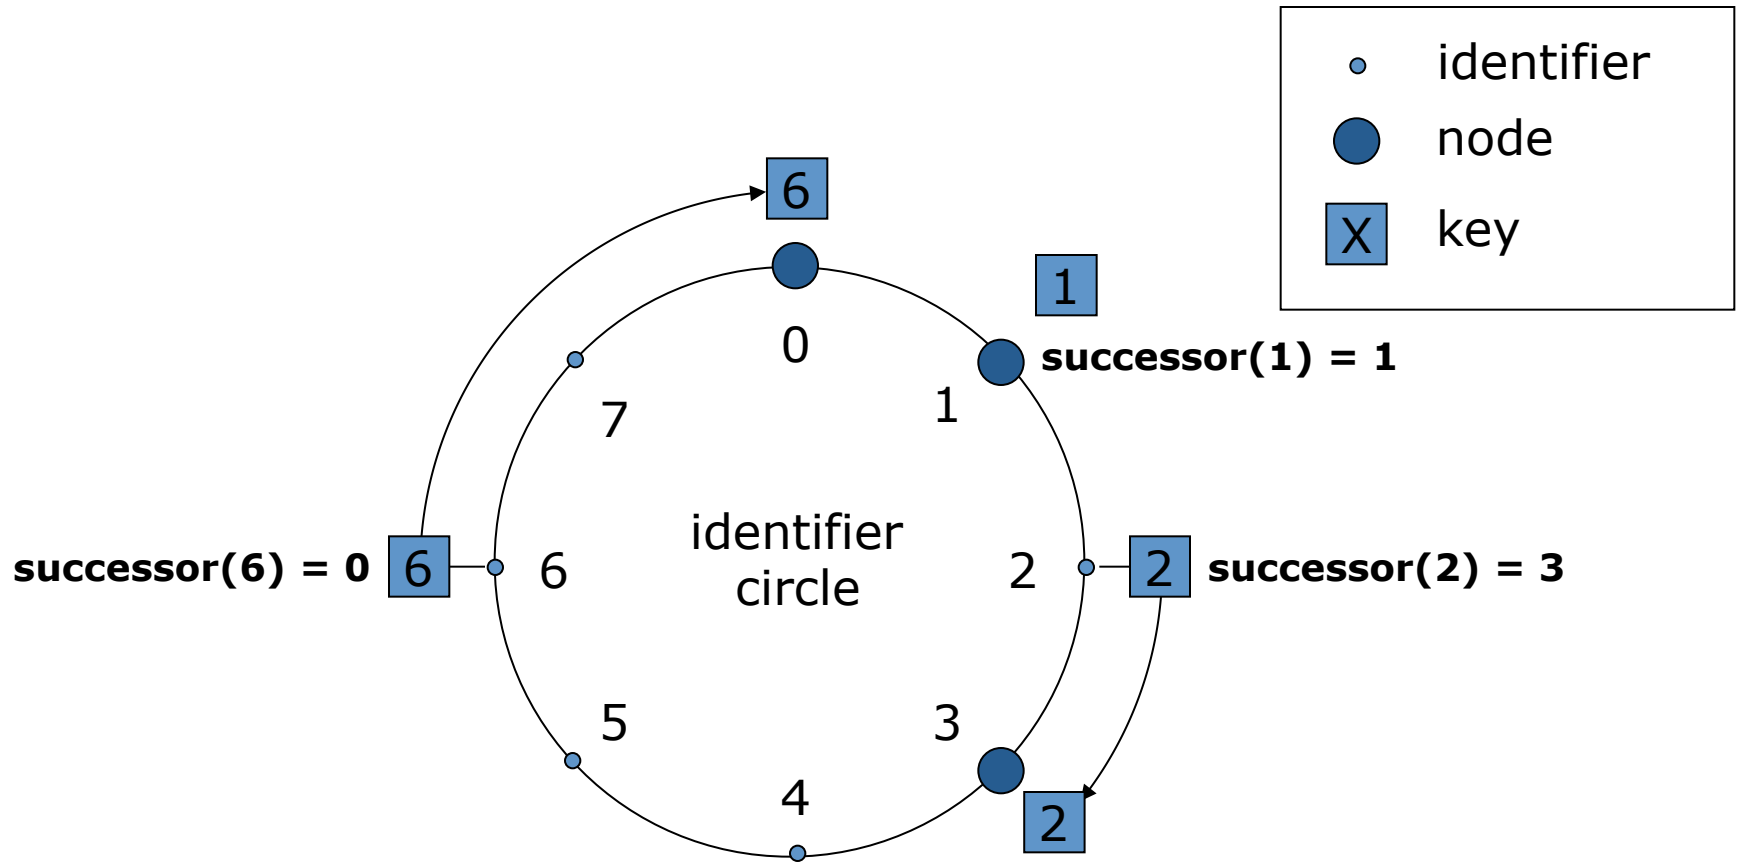
- Consistent hash function assigns each node and key an  $m$ -bit *identifier*.
- SHA-1 is used as a base hash function.
- A node's identifier is defined by hashing the node's IP address.
- A key identifier is produced by hashing the key (chord doesn't define this. Depends on the application).
  - $ID(\text{node}) = \text{hash}(\text{IP}, \text{Port})$
  - $ID(\text{key}) = \text{hash}(\text{key})$

# Consistent Hashing

- In an  $m$ -bit identifier space, there are  $2^m$  identifiers.
- Identifiers are ordered on an *identifier circle* modulo  $2^m$ .
- The identifier ring is called *Chord ring*.
- Key  $k$  is assigned to the first node whose identifier is equal to or follows (the identifier of)  $k$  in the identifier space.
- This node is the successor node of key  $k$ , denoted by *successor( $k$ )*.




# Consistent Hashing - Successor Nodes




# Consistent Hashing – Join and Departure

- When a node  $n$  joins the network, certain keys previously assigned to  $n$ 's successor now become assigned to  $n$ .
- When node  $n$  leaves the network, all of its assigned keys are reassigned to  $n$ 's successor.

# Consistent Hashing – Node Join



# Consistent Hashing – Node Dep.



# A Simple Key Lookup

- If each node knows only how to contact its current successor node on the identifier circle, all node can be visited in linear order.
- Queries for a given identifier could be passed around the circle via these successor pointers until they encounter the node that contains the key.

# A Simple Key Lookup

- Pseudo code for finding successor:

*// ask node n to find the successor of id*

`n.find_successor(id)`

**if** ( $id \in (n, \text{successor}]$ )

**return** *successor*;


**else**

*// forward the query around the circle*

**return** *successor.find\_successor(id)*;

# A Simple Key Lookup

- The path taken by a query from node 8 for key 54:



# Scalable Key Location

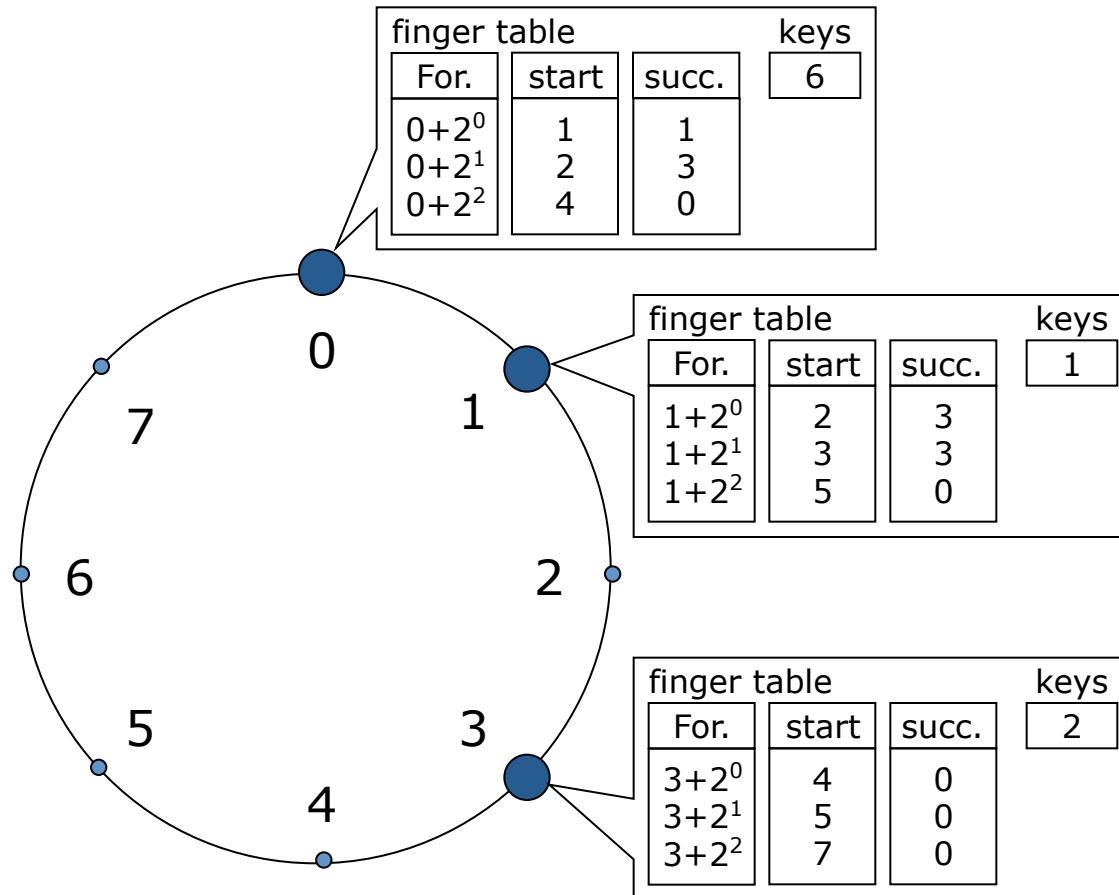
- To accelerate lookups, Chord maintains additional routing information.
- This additional information is not essential for correctness, which is achieved as long as each node knows its correct successor.



# Scalable Key Location – Finger Tables

- Each node  $n$  maintains a routing table with up to  $m$  entries (which is in fact the number of bits in identifiers), called *finger table*.
- The  $i^{\text{th}}$  entry in the table at node  $n$  contains the identity of the *first* node  $s$  that succeeds  $n$  by at least  $2^{i-1}$  on the identifier circle.
- $s = \text{successor}(n + 2^{i-1})$ .
- $s$  is called the  $i^{\text{th}}$  *finger* of node  $n$ , denoted by  $n.\text{finger}(i)$

# Scalable Key Location – Finger Tables




# Scalable Key Location – Finger Tables

- A finger table entry includes both the Chord identifier and the IP address (and port number) of the relevant node.
- The first finger of  $n$  is the immediate successor of  $n$  on the circle.

# Scalable Key Location – Example query

- The path a query for key 54 starting at node 8:



# Scalable Key Location – A characteristic

- Since each node has finger entries at power of two intervals around the identifier circle, each node can forward a query at least halfway along the remaining distance between the node and the target identifier. From this intuition follows a theorem:

*Theorem: With high probability, the number of nodes that must be contacted to find a successor in an  $N$ -node network is  $O(\log N)$ .*

# Node Joins and Stabilizations

- The most important thing is the successor pointer.
- If the successor pointer is ensured to be up to date, which is sufficient to guarantee correctness of lookups, then finger table can always be verified.
- Each node runs a “stabilization” protocol periodically in the background to update successor pointer and finger table.

# Node Joins and Stabilizations

- “Stabilization” protocol contains 6 functions:
  - create()
  - join()
  - stabilize()
  - notify()
  - fix\_fingers()
  - check\_predecessor()

# Node Joins – *join()*

- When node *n* first starts, it calls *n.join(n')*, where *n'* is any known Chord node.
- The *join()* function asks *n'* to find the immediate successor of *n*.
- *join()* does not make the rest of the network aware of *n*.



# Node Joins – join()

*// create a new Chord ring.*

**n.create()**

*predecessor = nil;*

*successor = n;*

*// join a Chord ring containing node n'.*

**n.join(n')**

*predecessor = nil;*

*successor = n'.find\_successor(n);*

# Node Joins – *stabilize()*

- Each time node  $n$  runs *stabilize()*, it asks its successor for its predecessor  $p$ , and decides whether  $p$  should be  $n$ 's successor instead.
- *stabilize()* notifies node  $n$ 's successor of  $n$ 's existence, giving the successor the chance to change its predecessor to  $n$ .
- The successor does this only if it knows of no closer predecessor than  $n$ .

# Node Joins – stabilize()

*// called periodically. verifies n's immediate  
// successor, and tells the successor about n.*

**n.stabilize()**

*x = successor.predecessor;*

*if ( $x \in (n, \text{successor})$ )*

*successor = x;*

*successor.notify(n);*


*// n' thinks it might be our predecessor.*

**n.notify(n')**

*if (predecessor is nil or  $n' \in (\text{predecessor}, n)$ )*

*predecessor = n';*

# Node Joins – Join and Stabilization



- **n joins**
  - predecessor = nil
  - n acquires  $n_s$  as successor via some  $n'$
- **n runs stabilize**
  - n notifies  $n_s$  being the new predecessor
  - $n_s$  acquires n as its predecessor
- **$n_p$  runs stabilize**
  - $n_p$  asks  $n_s$  for its predecessor (now n)
  - $n_p$  acquires n as its successor
  - $n_p$  notifies n
  - n will acquire  $n_p$  as its predecessor
- **all predecessor and successor pointers are now correct**
- **fingers still need to be fixed, but old fingers will still work**

# Node Joins – fix\_fingers()

- Each node periodically calls *fix fingers* to make sure its finger table entries are correct.
- It is how new nodes initialize their finger tables
- It is how existing nodes incorporate new nodes into their finger tables.

# Node Joins – fix\_fingers()

*// called periodically. refreshes finger table entries.*

**n.fix\_fingers()**

next = next + 1 ;

if (next > m)

next = 1 ;

finger[next] = find\_successor(n +  $2^{\text{next}-1}$ );

*// checks whether predecessor has failed.*

**n.check\_predecessor()**

if (predecessor has failed)

predecessor = nil;

# Node Failures

- Key step in failure recovery is maintaining correct successor pointers
- To help achieve this, each node maintains a *successor-list* of its  $r$  nearest successors on the ring. *Hence, all  $r$  successors would have to simultaneously fail in order to disrupt the Chord ring.*
- If node  $n$  notices that its successor has failed, it replaces it with the first live entry in the list
- Successor lists are stabilized as follows:
  - node  $n$  reconciles its list with its successor  $s$  by copying  $s$ 's successor list, removing its last entry, and prepending  $s$  to it.
  - If node  $n$  notices that its successor has failed, it replaces it with the first live entry in its successor list and reconciles its successor list with its new successor.

# Chord – The Math

- Each node maintains  $O(\log N)$  state information and lookups needs  $O(\log N)$  messages
- Every node is responsible for about  $K/N$  keys ( $N$  nodes,  $K$  keys)
- When a node joins or leaves an  $N$ -node network, only  $O(K/N)$  keys change hands (and only to and from joining or leaving node)



# Interesting Simulation Results

- Adding virtual nodes as an indirection layer can significantly improve load balance.
- The average path length is about  $\frac{1}{2}(\log N)$ .
- Maintaining a set of alternate nodes for each finger and route the queries by selecting the closest node according to network proximity metric improves routing latency effectively.
- Recursive lookup style is faster iterative style

# Applications: Time-shared storage

- for nodes with intermittent connectivity (server only occasionally available)
- Store others ' data while connected, in return having their data stored while disconnected
- Data 's name can be used to identify the live Chord node (content-based routing)

# Applications: Chord-based DNS

- DNS provides a lookup service

keys: host names values: IP addresses

Chord could hash each host name to a key

- Chord-based DNS:
  - no special root servers
  - no manual management of routing information
  - no naming structure
  - can find objects not tied to particular machines

# Summary

- Simple, powerful protocol
- Only operation: map a key to the responsible node
- Each node maintains information about  $O(\log N)$  other nodes
- Lookups via  $O(\log N)$  messages
- Scales well with number of nodes
- Continues to function correctly despite even major changes of the system

Thanks for your attention!