

Index Tuning Exercises

Database Management and Performance Tuning

The exercises use the following tables:

- Employee(ssnum, name, dept, manager, salary)
- Student(ssnum, name, course, grade, stipend, evaluation)

1. When the Student relation was created, a nonclustering index was created on **name**. However, the following query does not use that index:

```
SELECT *  
FROM Student  
WHERE name = 'Bayer'
```

Please explain any possible reason and provide a tuning solution.

2. You discover that the following important query is too slow. It is a simple query, and there is a nonclustering index on salary. You have tried to update the catalog statistics, but that did not help.

```
SELECT *  
FROM Employee  
WHERE salary/12 = 4000
```

Please explain any possible reason and provide a tuning solution.

3. The system uses the index on **salary** without improving performance of the following query:

```
SELECT *  
FROM Employee  
WHERE salary = 48000
```

Please explain any possible reason and provide a tuning solution.

4. Your system has pages with size of 2 KB. The Student records are very long (about 1 KB) because of the length of the **evaluation** attribute. There is a clustering index on **ssnum**, but the table suffers overflow chaining when new **evaluation** data is added. Please explain any possible reason and provide a tuning solution.

5. Suppose there are 30 Employee records per page. Each employee belongs to one of 50 departments. Should you put a nonclustering index on **dept** to speed up the following query?

```
SELECT ssnum  
FROM Employee  
WHERE dept = 'IT'
```

Please explain the reason.

6. Suppose there are 30 Employee records per pages. However, in this case, there are 5000 departments. Should you put a nonclustering index on **dept** to support the same query as in *Exercise 5*? Please explain the reason.

7. Auditors take a copy of the Employee file to which they wish to apply a statistic analysis. They allow no updates but want to support the following accesses:

- a) Count all the employees that have a certain salary (frequent).
- b) Find the employees that have the maximum (or minimum) salary within a particular department (frequent).
- c) Find the employee with a certain **ssnum** (rare).

Initially there is no index. Please design indexes to support the above queries.

8. Suppose that the student **stipend** correspond to monthly salaries, whereas the employee **salaries** are yearly. To find out which employees are paid as much as which students, we have two choices.

```
SELECT *  
FROM Employee, Student  
WHERE salary = 12*stipend
```

or

```
SELECT *  
FROM Employee, Student  
WHERE salary/12 = stipend
```

Which is better? Please explain the reason.

9. A purchasing department maintains the relation

`Onorder(supplier, part, quantity, price)`

The department makes the following queries to `Onorder`:

- a) Add a record, specifying all attributes (very frequent).
- b) Delete a record, specifying supplier and part (very frequent).
- c) Find the total quantity of a given part on order (frequent).
- d) Find the total value of the orders to a given supplier (rare).

Please design indexes to support the above queries.

10. A table has a clustering B^+ -tree index on `ssnum` and performs simple retrievals and updates of records based on `ssnum`. The performance is still not good. Please explain any possible reason and provide a tuning solution.

11. Ellis Island is a small island south of Manhattan through which flowed some 17 million immigrants to the United States between the late 1800s and the mid-1900s. Immigration workers filled in some 200 fields on each immigrants, containing information such as **last name**, **first name**, **city of origin**, **ship taken**, **nationality**, **religion**, **arrival date**, and so on. You are to design a database management system to allow the approximately 100 million descendants of these 17 million to retrieve the record of their ancestors.

To identify an immigrant, the querier must know the **last name** of the immigrant as well as some other information. Most queriers will know the **last name** and either the **first name** or the **year of arrival**. Please design indexes to support such queries.

12. An airline manages 1000 flights a day. In their database, there is a table for each flight (called Flight) containing a flight identifier, a seat identifier, and a passenger name. There is also a master table (called Totals) that has the total number of passengers in each flight. Each reservation transaction updates a particular Flight table and the Totals table. They have a performance bottleneck whose symptom is high lock contention. They try to resolve this by breaking each reservation transaction into two: one that updates the appropriate Flight table and the other that updates the Totals table. That helps but not enough. Please explain any possible reasons and provide a tuning solution.