

# Database Management and Performance Tuning

## Index Tuning

Pei Li

University of Zurich  
Institute of Informatics

Unit 4

Acknowledgements: The slides are provided by Nikolaus Augsten and adapted from “Database Tuning” by Dennis Shasha and Philippe Bonnet.

- 1 Index Tuning
  - Query Types
  - Index Types

# Query Types

- Different indexes are good for different query types.
- We identify categories of queries with different index requirements.

# Query Types

- **Point query**: returns at most one record

```
SELECT name  
FROM Employee  
WHERE ID = 8478
```

- **Multipoint query**: returns multiple records based on equality condition

```
SELECT name  
FROM Employee  
WHERE department = 'IT'
```

- **Range query** on  $X$  returns records with values in interval of  $X$

```
SELECT name  
FROM Employee  
WHERE salary >= 155000
```

# Query Types

- **Prefix match query**: given an ordered sequence of attributes, the query specifies a condition on a prefix of the attribute sequence
- **Example**: attribute sequence: lastname, firstname, city
  - The following are prefix match queries:
    - `lastname='Gates'`
    - `lastname='Gates' AND firstname='George'`
    - `lastname='Gates' AND firstname like 'Ge%'`
    - `lastname='Gates' AND firstname='George' AND city='San Diego'`
  - The following are **not** prefix match queries:
    - `firstname='George'`
    - `lastname LIKE '%ates'`

# Query Types

- **Extremal query**: returns records with max or min values on some attributes

```
SELECT name  
FROM Employee  
WHERE salary = MAX(SELECT salary FROM Employee)
```

- **Ordering query**: orders records by some attribute value

```
SELECT *  
FROM Employee  
ORDER BY salary
```

- **Grouping query**: partition records into groups; usually a function is applied on each partition

```
SELECT dept, AVG(salary)  
FROM Employee  
GROUP BY dept
```

# Query Types

- **Join queries:** link two or more tables

- **Equality join:**


```
SELECT Employee.ssnum  
FROM Employee, Student  
WHERE Employee.ssnum = Student.ssnum
```

- Join with **non-equality** condition:

```
SELECT e1.ssnum  
FROM Employee e1, Employee e2  
WHERE e1.manager = e2.ssnum  
AND e1.salary > e2.salary
```

# What is an Index?

- An **index** is a data structure that supports efficient access to data:



- Index tuning **essential** to performance!
- **Improper index selection** can lead to:
  - indexes that are maintained but never used
  - files that are scanned in order to return a single record
  - multitable joins that run for hours or days



# Key of an Index


- **Search key** or simply “key” of an index:
  - single attribute or sequence of attributes
  - values on key attributes used to access records in table
- **Sequential Key:**
  - value is monotonic with insertion order
  - examples: time stamp, counter
- **Non-sequential Key:**
  - value unrelated to insertion order
  - examples: social security number, last name
- **Note:** index key different from key in normalization theory
  - normalization theory: key attributes have unique values
  - index key: not necessarily unique

# Index Characteristics

- Indexes can often be viewed as **trees** ( $B^+$ -tree, hash)
  - some nodes are in main memory (e.g., root)
  - nodes deeper down in tree are less likely to be in main memory
- **Number of levels**: number of nodes in root-leaf path
  - a node is typically a disk block
  - one block read required per level
  - reading a block costs several milliseconds (involves disk seek)
- **Fanout**: number of children a node can have
  - large fanout means few levels
- **Overflow strategy**: insert into a full node  $n$ 
  - $B^+$ -tree: split  $n$  into  $n$  and  $n'$ , both at same distance from root
  - overflow chaining:  $n$  stores pointer to new node  $n'$

# Sparse vs. Dense

- **Sparse index:** pointers to disk pages
  - at most one pointer per disk page
  - usually much less pointers than records
- **Dense index:** pointers to individual records
  - one key per record
  - usually more keys than sparse index
  - optimization: store repeating keys only once, followed by pointers



# Sparse vs. Dense

- Number of pointers:

$\text{ptrs in dense index} = \text{records per page} \times \text{ptrs in sparse index}$

- Pro sparse: less pointers
  - typically record size is smaller than page size
  - less pointers result in less levels (and disk accesses)
  - uses less space
- Pro dense: index may “cover” query

# Covering Index

- **Covering index:**

- answers read query within index structure
- fast, since data is not accessed

- **Example 1:** dense index on lastname

```
SELECT COUNT(lastname) WHERE lastname='Smith'
```

- **Example 2:** dense index on A, B, C (in that order)

- covered query:

```
SELECT B, C  
FROM R  
WHERE A = 5
```

- covered query, but not prefix:


```
SELECT A, C  
FROM R  
WHERE B = 5
```

- non-covered query: D requires data access

```
SELECT B, D  
FROM R  
WHERE A = 5
```

# Clustering vs. Non-Clustering

- **Clustering index** on attribute  $X$   
(also *primary index*)
  - records are grouped by attribute  $X$  on disk
  - $B^+$ -tree: records sorted by attribute  $X$
  - only one clustering index per table
  - dense or sparse
- **Non-clustering index** on attribute  $X$   
(also *secondary index*)
  - no constraint on table organization
  - more than one index per table
  - always dense



# Clustering Indexes

- Can be **sparse**:
  - fewer pointers than non-clustering index (always dense!)
  - if record is small, save one disk access per record access
- Good for **multi-point queries**:
  - equality access on non-unique attribute
  - all result records are on consecutive pages
  - example: look up last name in phone book
- Good for **range, prefix, ordering** queries:
  - works if clustering index is implemented as  $B^+$ -tree
  - prefix example: look up all last names starting with 'St' in phone book
  - result records are on consecutive pages
- Good for **equality join**:
  - fast also for join on non-key attributes
  - index on one table: indexed nested-loop
  - index on both tables: merge-join
- **Overflow pages** reduce efficiency:
  - if disk page is full, overflowing records go to overflow pages
  - overflow pages require additional disk accesses

# Equality Join with Clustering Index

- Example query:

```
SELECT Employee.ssnum, Student.course  
FROM Employee, Student  
WHERE Employee.firstname = Student.firstname
```

- Index on Employee.firstname: use index nested loop join
  - for each student look up employees with same first name
  - all matching employees are on consecutive pages
- Index on both firstname attributes: use merge join
  - read both tables in sorted order and merge ( $B^+$ -tree)
  - each page read exactly once
  - works also for hash indexes with same hash function



# Clustering Index and Overflow Pages

- Why **overflow pages**?
  - clustering index stores records on consecutive disk pages
  - insertion between two consecutive pages not possible
  - if disk page is full, overflowing records go to overflow pages
- Additional disk access for overflow page: **reduced speed**
- Overflow pages can **result from**:
  - inserts
  - updates that change key value
  - updates that increase record size (e.g., replace NULL by string)
- **Reorganize** index:
  - invoke special tool
  - or simply drop and re-create index

# Overflow Strategies

- Tune free space in disk pages:
  - Oracle, DB2: pctfree (0 is full), SQLServer: fillfactor (100 is full)
  - free space in page is used for new or growing records
  - little free space: space efficient, reads are faster
  - much free space: reduced risk of overflows
- Overflow strategies:
  - split: split full page into two half-full pages and link new page  
e.g.,  $A \rightarrow B \rightarrow C$ , splitting  $B$  results in  $A \rightarrow B' \rightarrow B'' \rightarrow C$   
(SQLServer)
  - chaining: full page has pointer to overflow page (Oracle)
  - append: overflowing records of all pages are appended at the end of the table (DB2)


# Non-Clustering Index

- Always useful for point queries.
- Particularly good if index covers query.
- Critical tables: covering index on all relevant attribute combinations
- Multi-point query (not covered): only if not too selective
  - $nR$ : number of records returned by query
  - $nP$ : number of disk pages in table
  - the  $nR$  records are uniformly distributed over all pages
  - thus query will read  $\min(nR, nP)$  disk pages
- Index may slow down highly selective multi-point query:
  - scan is by factor 2–10 faster than accessing all pages with index
  - thus  $nR$  should be significantly smaller than  $nP$

# Non-Clustering Index and Multi-point Queries – Example

- Example 1:
  - records size:  $50B$
  - page size:  $4kB$
  - attribute  $A$  takes 20 different values (evenly distributed among records)
  - does non-clustering index on  $A$  help?
- Evaluation:
  - $nR = m/20$  ( $m$  is the total number of records)
  - $nP = m/80$  (80 records per page)
  - $m/20 > m/80$  thus index does not help
- Example 2: as above, but record size is  $2kB$
- Evaluation:
  - $nR = m/20$  ( $m$  is the total number of records)
  - $nP = m/2$  (2 records per page)
  - $m/20 \ll m/2$  thus index might be useful

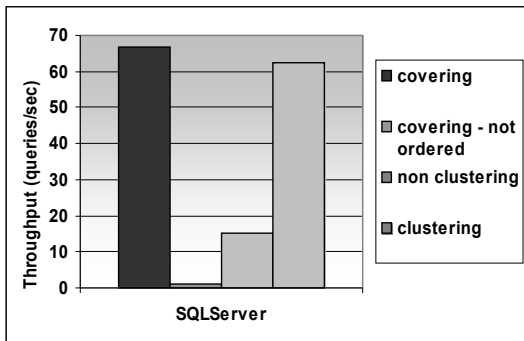
# Clustering vs. Non-Clustering Index



- multi-point query with selectivity 100/1M records (0.01%)
- clustering index much faster than non-clustering index
- full table scan (no index) orders of magnitude slower than index

DB2 UDB V7.1, Oracle 8.1, SQL Server 7 on Windows 2000


# Covering vs. Non-Covering Index



- prefix match query on sequence of attributes
- covering: index covers query, query condition on prefix
- covering, not ordered: index covers query, but condition not prefix
- non-clustering: non-covering index, query condition on prefix
- clustering: sparse index, query condition on prefix


SQL Server 7 on Windows 2000

# Non-Covering vs. Table Scan



- query: range query
- non clustering: non-clustering non-covering index
- scan: no index, i.e., table scan required
- index is faster if less than 15% of the records are selected


# Index Maintenance - DB2



- query: batch of 100 multi-point queries, pctfree=0 (data pages full)
- performance degrades with insertion
- overflow records simply appended
- query traverses index and then scans all overflow records
- reorganization helps




# Index Maintenance - SQL Server



- fillfactor=100 (data pages full)
- performance degrades with insertion
- overflow chain maintained for overflowing page
- extra disk access
- reorganization helps

# Index Maintenance - Oracle



- `pctfree = 0` (data pages full), performance degrades with insertion
- all indexes in Oracle are non-clustering
- index-organized table is clustered by primary key
- recreating index does not reorganize table
- maintenance: export and re-import table to reorganize :-)