

# Database Management and Performance Tuning

## Concurrency Tuning

Pei Li

University of Zurich  
Institute of Informatics

Unit 7

Acknowledgements: The slides are provided by Nikolaus Augsten and adapted from “Database Tuning” by Dennis Shasha and Philippe Bonnet.

# Outline

- 1 Index Tuning
  - Index Tuning Examples
- 2 Concurrency Tuning
  - Introduction to Transactions

# Index Tuning Examples

- The examples use the following tables:
  - `Employee(ssnum,name,dept,manager,salary)`
  - `Student(ssnum,name,course,grade,stipend,evaluation)`

# Exercise 1 – Query for Student by Name

- Student was created with non-clustering index on name.

- Query:

```
SELECT *  
FROM Student  
WHERE name='Bayer'
```

- Problem: Query does not use index on name.
- Solution: Try updating the catalog statistics.
  - Oracle, Postgres: ANALYZE
  - SQL Server: sp\_createstats
  - DB2: RUNSTATS

## Exercise 2 – Query for Salary I

- Non-clustering index on salary.
- Catalog statistics are up-to-date.

- Query:

```
SELECT *  
FROM Employee  
WHERE salary/12 = 4000
```

- Problem: Query is too slow.
- Solution: Index not used because of the arithmetic expression.  
Two Options:

- Rewrite query:

```
SELECT *  
FROM Employee  
WHERE salary = 48000
```

- Use function based index.

## Exercise 3 – Query for Salary II

- Non-clustering index on salary.
- Catalog statistics are up-to-date.
- Query:

```
SELECT *  
FROM Employee  
WHERE salary = 48000
```

- **Problem:** Query still does not use index. What could be the reason?
- **Solution:** The index is non-clustering. Many employees have a salary of 48000, thus the index may not help. It may still help for other, less frequent, salaries!

## Exercise 4 – Clustering Index and Overflows

- Clustering index on `Student.ssnum`
- Page size: *2kB*
- Record size in `Student` table: *1KB* (evaluation is a long text)
- **Problem:** Overflow when new evaluations are added.
- **Solution:** Clustering index does not help much due to large record size. A non-clustering index avoids overflows.

## Exercise 5 – Non-clustering Index I

- Employee table:
  - 30 employee records per page
  - each employee belongs to one of 50 departments (dept)
  - the departments are of similar size
- Query:

```
SELECT ssnnum
FROM Empl_yee
WHERE dept = 'IT'
```
- Problem: Does a non-clustering index on Employee.dept help?
- Solution: Only if the index covers the query.
  - $30/50=60\%$  of the pages will have a record with dept = 'IT'
  - table scan is faster than accessing 3/5 of the pages in random order



## Exercise 6 – Non-clustering Index II

- Employee table:
  - 30 employee records per page
  - each employee belongs to one of 5000 departments (dept)
  - the departments are of similar size
- Query:

```
SELECT ssnnum
FROM Empl_yee
WHERE dept = 'IT'
```
- Problem: Does a non-clustering index on Employee.dept help?
- Solution: Only if the index covers the query.
  - only  $30/5000=0.06\%$  of the pages will have a record with dept='IT'
  - table scan is slower

## Exercise 7 – Statistical Analysis

- Auditors run a **statistical analysis** on a copy of Employee.
- **Queries:**
  - count employees with a certain salary (frequent)
  - find employees with maximum or minimum salary within a particular department (frequent)
  - find an employee by its social security number (rare)
- **Problem:** Which indexes to create?
- **Solution:**
  - non-clustering index on salary (covers the query)
  - clustering composite index on (dept, salary) using a  $B^+$ -tree (all employees with the maximum salary are on consecutive pages)
  - non-clustering hash index on ssnum

## Exercise 8 – Algebraic Expressions

- Student stipends are monthly, employee salaries are yearly.
- **Query:** Which employee is paid as much as which student?
- There are **two options** to write the query:

```
SELECT *  
FROM Employee, Student  
WHERE salary = 12*stipend
```

```
SELECT *  
FROM Employee, Student  
WHERE salary/12 = stipend
```

- Index on a table with an algebraic expression not used.
- **Problem:** Which query is better?

## Exercise 8 – Solution

- If index on only **one table**, it should be used.
- Index on **both tables**, **clustering** on larger table: use it.
- Index on **both tables**, **non-clustering** on larger table:
  - small table has (much) less records than large table has pages:  
use index on large table
  - otherwise: use index on small table

## Exercise 9 – Purchasing Department

- Purchasing department maintains table `Onorder(supplier,part,quantity,price)`.
- The table is heavily used during the opening hours, but not over night.
- **Queries:**
  - Q1: add a record, all fields specified (very frequent)
  - Q2: delete a record, `supplier` and `part` specified (very frequent)
  - Q3: find total quantity of a given part on order (frequent)
  - Q4: find the total value on order to a given supplier (rare)
- **Problem:** Which indexes should be used?

## Exercise 9 – Solution

- **Queries:**
  - Q1: add a record, all fields specified (very frequent)
  - Q2: delete a record, `supplier` and `part` specified (very frequent)
  - Q3: find total quantity of a given part on order (frequent)
  - Q4: find the total value on order to a given supplier (rare)
- **Solution:** Clustering composite  $B^+$ -tree index on `(part, supplier)`.
  - eliminate overflows over night
  - attribute order important to support query Q3
  - hash index will not work for query Q3 (prefix match query)
- **Discussion:** Non-clustering index on `supplier` to answer query Q4?
  - index must be maintained and will hurt the performance of much more frequent queries Q1 and Q2
  - index does not help much if there are only few different suppliers

## Exercise 10 – Point Query Too Slow

- Employee has a clustering  $B^+$ -tree index on `ssnum`.
- Queries:
  - retrieve employee by social security number (`ssnum`)
  - update employee with a specific social security number
- Problem: Throughput is still not enough.
- Solution: Use hash index instead of  $B^+$ -tree (faster for point queries).

# Exercise 11 – Historical Immigrants Database

- Digitalized database of **US immigrants** between 1800 and 1900:
  - 17M records
  - each record has approx. 200 fields  
e.g., last name, first name, city of origin, ship taken, etc.
- **Queries** retrieve immigrants:
  - by last name and at least one other attribute
  - second attribute is often first name (most frequent) or year
- **Problem**: Efficiently serve 2M descendants of the immigrants. . .



# Exercise 11 – Solution

- Clustering  $B^+$ -tree index on (lastname,firstname):
  - no overflow since database does not have updates
  - use high fill factor to increase space utilization
  - key compression should be used (long key, no update)
  - index useful also for prefix queries on lastname
- Composite non-clustering index on (lastname,year):
  - no maintenance cost (no updates)
  - attributes probably selective enough
- Non-clustering indexes on all frequent attribute combinations?
  - no maintenance, thus only limitation is space overhead
  - useful only if selective enough

## Exercise 12 – Flight Reservation System

- An airline manages 1000 flights and uses the tables:
  - `Flight(flightID, seatID, passanger-name)`
  - `Totals(flightID, number-of-passangers)`
- **Query:** Each reservation
  - adds a record to `Flight`
  - increments `Totals.number-of-passangers`
- Queries are **separate transactions**.
- **Problem:** Lock contention on `Totals`.
- **Solution:**
  - `Totals` is a small table (1000 small records) and fits on few pages.
  - Without index, update scans table and scanned records are locked.
  - Clustering index on `flightID` avoids table scan and thus lock contention (row locking assumed).

# Outline

- 1 Index Tuning
  - Index Tuning Examples
- 2 Concurrency Tuning
  - Introduction to Transactions

# What is a Transaction?<sup>1</sup>

- A **transaction** is a unit of program execution that accesses and possibly updates various data items.
- **Example:** transfer \$50 from account  $A$  to account  $B$ 
  1.  $R(A)$
  2.  $A \leftarrow A - 50$
  3.  $W(A)$
  4.  $R(B)$
  5.  $B \leftarrow B + 50$
  6.  $W(B)$
- Two **main issues**:
  1. concurrent execution of multiple transactions
  2. failures of various kind (e.g., hardware failure, system crash)

<sup>1</sup> Slides of section “Introduction to Transactions” are adapted from the slides “Database System Concepts”, 6<sup>th</sup> Ed., Silberschatz, Korth, and Sudarshan

# ACID Properties

- Database system must guarantee **ACID for transactions**:
  - **Atomicity**: either all operations of the transaction are executed or none
  - **Consistency**: execution of a transaction in isolation preserves the consistency of the database
  - **Isolation**: although multiple transactions may execute concurrently, each transaction must be unaware of the other concurrent transactions.
  - **Durability**: After a transaction completes successfully, changes to the database persist even in case of system failure.

# Atomicity

- **Example:** transfer \$50 from account  $A$  to account  $B$ 
  1.  $R(A)$
  2.  $A \leftarrow A - 50$
  3.  $W(A)$
  4.  $R(B)$
  5.  $B \leftarrow B + 50$
  6.  $W(B)$
- What if **failure** (hardware or software) after step 3?
  - money is lost
  - database is inconsistent
- **Atomicity:**
  - either all operations or none
  - updates of partially executed transactions not reflected in database

# Consistency

- **Example:** transfer \$50 from account  $A$  to account  $B$ 
  1.  $R(A)$
  2.  $A \leftarrow A - 50$
  3.  $W(A)$
  4.  $R(B)$
  5.  $B \leftarrow B + 50$
  6.  $W(B)$
- **Consistency in example:** sum  $A + B$  must be unchanged
- **Consistency in general:**
  - explicit integrity constraints (e.g., foreign key)
  - implicit integrity constraints (e.g., sum of all account balances of a bank branch must be equal to branch balance)
- **Transaction:**
  - must see consistent database
  - during transaction inconsistent state allowed
  - after completion database must be consistent again

# Isolation – Motivating Example

- **Example:** transfer \$50 from account  $A$  to account  $B$ 
  1.  $R(A)$
  2.  $A \leftarrow A - 50$
  3.  $W(A)$
  4.  $R(B)$
  5.  $B \leftarrow B + 50$
  6.  $W(B)$
- Imagine second transaction  $T_2$ :
  - $T_2 : R(A), R(B), \text{print}(A + B)$
  - $T_2$  is executed between steps 3 and 4
  - $T_2$  sees an inconsistent database and gives wrong result



# Isolation

- **Trivial isolation**: run transactions serially
- **Isolation** for concurrent transactions: For every pair of transactions  $T_i$  and  $T_j$ , it appears to  $T_i$  as if either  $T_j$  finished execution before  $T_i$  started or  $T_j$  started execution after  $T_i$  finished.
- **Schedule**:
  - specifies the **chronological order** of a sequence of instructions from various transactions
  - **equivalent schedules** result in identical databases if they start with identical databases
- **Serializable** schedule:
  - equivalent to some serial schedule
  - serializable schedule of  $T_1$  and  $T_2$  is either equivalent to  $T_1, T_2$  or  $T_2, T_1$

# Durability

- When a transaction is done it **commits**.
- **Example**: transaction commits too early
  - transaction writes *A*, then commits
  - *A* is written to the disk buffer
  - then system crashes
  - value of *A* is lost
- **Durability**: After a transaction has committed, the changes to the database persist even in case of system failure.
- **Commit** only after all changes are permanent:
  - either written to log file or directly to database
  - database must recover in case of a crash

# Locks

- A **lock** is a mechanism to **control concurrency** on a data item.
- Two types of locks on a data item  $A$ :
  - **exclusive** –  $xL(A)$ : data item  $A$  can be both read and written
  - **shared** –  $sL(A)$ : data item  $A$  can only be read.
- **Lock request** are made to concurrency control manager.
- Transaction is **blocked** until lock is granted.
- **Unlock  $A$**  –  $uL(A)$ : release the lock on a data item  $A$

# Lock Compatibility

- Lock **compatibility matrix**:

$T_1 \downarrow T_2 \rightarrow$	shared	exclusive
shared	true	false
exclusive	false	false

- $T_1$  holds **shared lock** on  $A$ :
  - shared lock is granted to  $T_2$
  - exclusive lock is not granted to  $T_2$
- $T_2$  holds **exclusive lock** on  $A$ :
  - shared lock is not granted to  $T_2$
  - exclusive lock is not granted to  $T_2$
- Shared locks can be shared by **any number** of transactions.

# Locking Protocol

- Example transaction  $T_2$  with locking:
  1.  $sL(A), R(A), uL(A)$
  2.  $sL(B), R(B), uL(B)$
  3.  $print(A + B)$
- $T_2$  uses locking, but is not serializable
  - $A$  and/or  $B$  could be updated between steps 1 and 2
  - printed sum may be wrong
- Locking protocol:
  - set of rules followed by all transactions while requesting/releasing locks
  - locking protocol restricts the set of possible schedules

# Pitfalls of Locking Protocols – Deadlock

- **Example:** two concurrent money transfers
  - $T_1$ :  $R(A)$ ,  $A \leftarrow A + 10$ ,  $R(B)$ ,  $B \leftarrow B - 10$ ,  $W(A)$ ,  $W(B)$
  - $T_2$ :  $R(B)$ ,  $B \leftarrow B + 50$ ,  $R(A)$ ,  $A \leftarrow A - 50$ ,  $W(A)$ ,  $W(B)$
  - possible concurrent scenario with locks:  
 $T_1.xL(A)$ ,  $T_1.R(A)$ ,  $T_2.xL(B)$ ,  $T_2.R(B)$ ,  $T_2.xL(A)$ ,  $T_1.xL(B)$ , ...
  - $T_1$  and  $T_2$  block each other – no progress possible
- **Deadlock:** situation when transactions block each other
- **Handling** deadlocks:
  - one of the transactions must be rolled back (i.e., undone)
  - rolled back transaction releases locks

# Pitfalls of Locking Protocols – Starvation

- **Starvation:** transaction continues to wait for lock
- **Examples:**
  - the same transaction is repeatedly rolled back due to deadlocks
  - a transaction continues to wait for an exclusive lock on an item while a sequence of other transactions are granted shared locks
- Well-designed concurrency manager **avoids starvation**.

# Two-Phase Locking

- Protocol that **guarantees serializability**.
- **Phase 1**: growing phase
  - transaction may obtain locks
  - transaction may not release locks
- **Phase 2**: shrinking phase
  - transaction may release locks
  - transaction may not obtain locks



# Two-Phase Locking – Example

- **Example:** two concurrent money transfers
  - $T_1$ :  $R(A), A \leftarrow A + 10, R(B), B \leftarrow B - 10, W(A), W(B)$
  - $T_2$ :  $R(A), A \leftarrow A - 50, R(B), B \leftarrow B + 50, W(A), W(B)$
- Possible **two-phase locking schedule**:
  1.  $T_1$  :  $xL(A), xL(B), R(A), R(B), W(A \leftarrow A + 10), uL(A)$
  2.  $T_2$  :  $xL(A), R(A), xL(B)$  (*wait*)
  3.  $T_1$  :  $W(B \leftarrow B - 10), uL(B)$
  4.  $T_2$  :  $R(B), W(A \leftarrow A - 50), W(B \leftarrow B + 50), uL(A), uL(B)$
- **Equivalent serial** schedule:  $T_1, T_2$