

Database Management and Performance Tuning

Concurrency Tuning

Pei Li

University of Zurich
Institute of Informatics

Unit 9

Acknowledgements: The slides are provided by Nikolaus Augsten and adapted from “Database Tuning” by Dennis Shasha and Philippe Bonnet.

Outline

- 1 Concurrency Tuning
 - Weaken Isolation Guarantees

Undesirable Phenomena of Concurrent Transactions

- Dirty read

- transaction reads data written by concurrent uncommitted transaction
- problem: read may return a value that was never in the database because the writing transaction aborted

- Non-repeatable read

- different reads on the same item within a single transaction give different results (caused by other transactions)
- e.g., concurrent transactions $T_1: x = R(A), y = R(A), z = y - x$ and $T_2: W(A = 2 * A)$, then z can be either zero or the initial value of A (should be zero!)

- Phantom read

- repeating the same query later in the transaction gives a different set of result tuples
- other transactions can insert new tuples during a scan
- e.g., “Q: get accounts with *balance* > 1000” gives two tuples the first time, then a new account with *balance* > 1000 is inserted by an other transaction; the second time Q gives three tuples

Isolation Guarantees (SQL Standard)

- **Read uncommitted**: dirty, non-repeatable, phantom
 - read locks released after read; write locks downgraded to read locks after write, downgraded locks released according to 2-phase locking
 - reads may access uncommitted data
 - writes do not overwrite uncommitted data
- **Read committed**: non-repeatable, phantom
 - read locks released after read, write locks according to 2-phase locking
 - reads can access only committed data
 - **cursor stability**: in addition, read is repeatable within single SELECT
- **Repeatable read**: phantom
 - 2-phase locking, but no range locks
 - phantom reads possible
- **Serializable**:
 - none of the undesired phenomenas can happen
 - enforced by 2-phase locking with range locks

Experiment: Read Commit vs. Serializable

- **Experimental setup:**


- T1: summation query: `SELECT SUM(balance) FROM Accounts`
- T2: money transfers between accounts
- row level locking

- **Parameter:** number of concurrent threads

- **Measure:**


- percentage of correct answers (over multiple tries)
- measure throughput

Experiment: Read Commit vs. Serializable



- **Read committed** allows sum of account balances after debit operation has taken place but before corresponding credit operation is performed – incorrect sum!

Experiment: Read Commit vs. Serializable



- Read committed: faster, but incorrect answers
- Serializable: always correct, but lower throughput

When To Weaken Isolation Guarantees?

- Query does not need exact answer (e.g., statistical queries)
 - example: count all accounts with balance > \$1000.
 - read committed is enough!
- Transactions with human interaction
 - example: flight reservation system
 - price for serializability too high!

Example: Flight Reservation System

- Reservation involves **three steps**:
 1. retrieve list of available seats
 2. let customer decide
 3. secure seat
- **Single transaction**:
 - seats are locked while customer decides
 - all other customers are blocked!
- **Two transactions**: (1) retrieve list, (2) secure seat
 - seat might already be taken when customer wants to secure it
 - more tolerable than blocking all other customers

Snapshot Isolation for Long Reads – The Problem

- Consider the following **scenario in a bank**:
 - read-only query Q : `SELECT SUM(deposit) FROM Accounts`
 - update transaction T : money transfer between customers A and B
- **2-Phase locking inefficient** for long read-only queries:
 - read-only queries hold lock on all read items
 - in our example, T must wait for Q to finish (Q blocks T)
 - deadlocks might occur:
 $T.xL(A)$, $Q.sL(B)$, $Q.sL(A)$ - wait, $T.xL(B)$ - wait
- **Read-committed** may lead to **incorrect** results:
 - Before transactions: $A = 50, B = 30$
 - $Q : sL(A), R(A) = 50, uL(A)$
 - $T : xL(A), xL(B), W(A \leftarrow A + 20), W(B \leftarrow B - 20), uL(A), uL(B)$
 - $Q : sL(B), R(B) = 10, uL(B)$
 - sum computed by Q for $A + B$ is 60 (instead of 80)

Snapshot Isolation for Long Reads

- **Snapshot isolation:** correct read-only queries without locking
 - read-only query Q with snapshot isolation
 - remember old values of all data items that change after Q starts
 - Q sees the values of the data items when Q started
- **Example:** bank scenario with snapshot isolation
 - Before transactions: $A = 50, B = 30$
 - $Q : R(A) = 50$
 - $T : xL(A), xL(B), W(A \leftarrow A + 20), W(B \leftarrow B - 20), uL(A), uL(B)$
 - $Q : R(B) = 30$ (read old value)
 - sum computed by Q for $A + B$ is 80 as it should be

Concurrency in Oracle

- “Read committed” in Oracle means:
 - non-repeatable and phantom reads are possible at the transaction level, but not within a single SQL statement
 - update conflict: if row is already updated, wait for updating transaction to commit, then update new row version (or ignore row if deleted) – no rollback!
 - possibly inconsistent state: transaction sees updates of other transaction only on the rows that itself updates
- “Serializable” in Oracle means:
 - phenomena: none of the three undesired phenomena can happen
 - update conflict: if two transactions update the same item, the transaction that updates it later must abort – rollback!
 - not serializable: snapshot isolation does not guarantee full serializability (skew writes)
- Similar in PostgreSQL.

Skew Writes: Snapshot Isolation Not Serializable

- **Example:** $A = 3, B = 17$
 - $T1 : A \leftarrow B$
 - $T2 : B \leftarrow A$
- **Serial execution:**
 - order $T1, T2$: $A = B = 17$
 - order $T2, T1$: $A = B = 3$
- **Snapshot isolation:**
 - $T1 : R(B) = 17$
 - $T2 : R(A) = 3$
 - $T1 : W(A \leftarrow 17)$
 - $T2 : W(B \leftarrow 3)$
 - result: $A = 17, B = 3$ (different from serial execution)


Snapshot Isolation

- **Advantages:** (assuming “serializable” of Oracle)
 - readers do not block writers (as with locking)
 - writers do not block readers (as with locking)
 - writers block writers only if they update the same row
 - performance similar to read committed
 - no dirty, non-repeatable, or phantom reads
- **Disadvantages:**
 - system must write and hold old versions of modified data (only data modified between start and end of read-only transaction)
 - does **not guarantee serializability** for read/write transactions
- **Implementation example:** Oracle 9i
 - no overhead: leverages before-image in rollback segment
 - expiration time of before-images configurable, “snapshot too old” failure if this value is too small

Snapshot Isolation – Summary


- Considerable **performance advantages** since reads are never blocked and do not block other transactions.
- **Not fully serializable**, although no dirty, non-repeatable, or phantom reads.

Experiment: Read Commit vs. Serializable



- Summation query with concurrent transfers between bank accounts.
- Oracle snapshot isolation: read-only summation query is not disturbed by concurrent transfer queries
- Summation (read-only) queries always give exact answer.

Experiment: Read Commit vs. Serializable



- Both “read commit” and “serializable” use snapshot isolation.
- “Serializable” rolls back transactions in case of write conflict.
- Summation queries always give exact answer.