

Database Management and Performance Tuning

Concurrency Tuning

Pei Li

University of Zurich
Institute of Informatics

Unit 10

Acknowledgements: The slides are provided by Nikolaus Augsten and adapted from “Database Tuning” by Dennis Shasha and Philippe Bonnet.

Outline

- 1 Concurrency Tuning
 - Transaction Chopping

Chopping Long Transactions

- Shorter transactions

- request less locks (thus they are less likely to be blocked or block an other transaction)
- require other transactions to wait less for a lock
- are better for logging

- Transaction chopping:

- split long transactions into short ones
- don't scarify correctness

Terminology

- **Transaction**: sequence of disc accesses (read/write)
- **Piece** of transaction: consecutive subsequence of database access.
 - example transaction $T : R(A), R(B), W(A)$
 - $R(A)$ and $R(A), R(B)$ are pieces of T
 - $R(A), W(A)$ is not a piece of T (not consecutive)
- **Chopping**: partitioning transaction it into pieces.
 - example transaction $T : R(A), R(B), W(A)$
 - $T_1 : R(A), R(B)$ and $T_2 : W(A)$ is a chopping of T

Split Long Transactions – Example 1

- **Bank** with accounts and branches:
 - each account is assigned to exactly one branch
 - branch balance is sum of accounts in that branch
 - customers can take out cash during day
- **Transactions** over night:
 - **update transaction**: reflect daily withdrawals in database
 - **balance checks**: customers ask for account balance (read-only)
- **Update transaction** T_{blob}
 - updates all account balances to reflect daily withdrawals
 - updates the respective branch balances
- **Problem**: balance checks are blocked by T_{blob} and take too long

Split Long Transactions – Example 1

- **Solution:** split update transactions T_{blob} into many small transactions
- **Variant 1:** each account update is one transaction which
 - updates one account
 - updates the respective branch balance
- **Variant 2:** each account update consists of two transactions
 - T_1 : update account
 - T_2 : update branch balance
- **Note:** isolation does not imply consistency
 - both variants maintain serializability (isolation)
 - variant 2: consistency (sum of accounts equal branch balance) compromised if only one of T_1 or T_2 commits.

Split Long Transactions – Example 2

- Bank scenario as in Example 1.
- Transactions:
 - update transaction: each transaction updates one account and the respective branch balance (variant 1 in Example 1)
 - balance checks: customers ask for account balance (read-only)
 - consistency (T'): compute account sum for each branch and compare to branch balance
- Splitting: T' can be split into transactions for each individual branch
- Serializability maintained:
 - consistency checks on different branches share no data item
 - updates leave database in consistent state for T'
- Note: update transaction can not be further split (variant 2)!
- Lessons learned:
 - sometimes transactions can be split without sacrificing serializability
 - adding new transaction to setting may invalidate all previous chopping

Formal Chopping Approach

- **Assumptions**: when can the chopping be applied?
- **Execution rules**: how must chopped transactions be executed?
- **Chopping graph**: which chopping is correct?

Assumptions for Transaction Chopping

1. **Transactions:** All transactions that run in an interval are known.
2. **Rollbacks:** It is known where in the transaction rollbacks are called.
3. **Failure:** In case of failure it is possible to determine which transactions completed and which did not.
4. **Variables:** The transaction code that modifies a program variable x must be reentrant, i.e., if the transaction aborts due to a concurrency conflict and then executes properly, x is left in a consistent state.

Execution Rules

1. **Execution order:** The execution of pieces obeys the order given by the transaction.
2. **Lock conflict:** If a piece is aborted due to a lock conflict, then it will be resubmitted until it commits.
3. **Rollback:** If a piece is aborted due to a rollback, then no other piece for that transaction will be executed.

The Transaction Chopping Problem

- **Given:** Set $A = \{T_1, T_2, \dots, T_n\}$ of (possibly) concurrent transactions.
- **Goal:** Find a chopping B of the transactions in A such that any serializable execution of the transactions in B (following the execution rules) is equivalent to some serial execution of the transactions in A . Such a chopping is said to be **correct**.
- **Note:** The “serializable” execution of B may be concurrent, following a protocol for serializability.

Chopping Graph

- We represent a specific chopping of transactions as a graph.
- **Chopping graph**: undirected graph with two types of edges.
 - nodes: each piece in the chopping is a node
 - C-edges: edge between any two conflicting pieces
 - S-edges: edge between any two sibling pieces
- **Conflicting pieces**: two pieces p and p' conflict iff
 - p and p' are pieces of different original transactions
 - both p and p' access a data item x and at least one modifies it
- **Sibling pieces**: two pieces p and p' are siblings iff
 - p and p' are neighboring pieces of the same original transactions

Chopping Graph – Example

- **Notation:** chopping of possibly concurrent transactions.
 - original transactions are denoted as T_1, T_2, \dots
 - chopping T_i results in pieces T_{i1}, T_{i2}, \dots
- **Example transactions:** ($T_1 : R(x), R(y), W(y)$ is split into T_{11}, T_{12})
 - $T_{11} : R(x)$
 - $T_{12} : R(y), W(y)$
 - $T_2 : R(x), W(x)$
 - $T_3 : R(y), W(y)$
- **Conflict edge** between nodes
 - T_{11} and T_2 (conflict on x)
 - T_{12} and T_3 (conflict on y)
- **Sibling edge** between nodes
 - T_{11} and T_{12} (same original transaction T_1)

Rollback Safe

- **Motivation:** Transaction T is chopped into T_1 and T_2 .
 - T_1 executes and commits
 - T_2 contains a rollback statement and rolls back
 - T_1 is already committed and will not roll back
 - in original transaction T rollback would also undo effect of piece T_1 !
- A chopping of transaction T is **rollback safe** if
 - T has no rollback statements or
 - all rollback statements are in the first piece of the chopping

Correct Chopping

Theorem (Correct Chopping)

A chopping is correct if it is rollback save and its chopping graph contains no SC-cycles.

- Chopping of previous example is correct (no SC-cycles, no rollbacks)
- If a chopping is not correct, then any further chopping of any of the transactions will not render it correct.
- If two pieces of transaction T are in an SC-cycle as a result of chopping T , then they will be in a cycle even if no other transactions (different from T) are chopped.

Private Chopping

- **Private chopping:** Given transactions T_1, T_2, \dots, T_n .
 $T_{i1}, T_{i2}, \dots, T_{ik}$ is a private chopping of T_i if
 - there is no SC-cycle in the graph with the nodes $\{T_1, \dots, T_{i1}, \dots, T_{ik}, \dots, T_n\}$
 - T_i is rollback save
- **Private chopping rule:** The chopping that consists of $private(T_1), private(T_2), \dots, private(T_n)$ is correct.
- **Implication:**
 - each transaction T_i can be chopped in isolation, resulting in $private(T_i)$
 - overall chopping is union of private choppings

Chopping Algorithm

1. Draw an S-edge between the R/W operations of a single transaction.
2. For each data item x produce a write list, i.e., a list of transactions that write this data item.
3. For each $R(x)$ or $W(x)$ in all transactions:
 - (a) look up the conflicting transactions in the write list of x
 - (b) draw a C-edge to the respective conflicting operations
4. Remove all S-edges that are involved in an SC-cycle.

Chopping Algorithm – Example

- **Transactions:** ($R_x = R(x)$, $W_x = W(x)$)
 - $T_1 : R_x, W_x, R_y, W_y$
 - $T_2 : R_x, W_x$
 - $T_3 : R_y, R_z, W_y$
- **Write lists:** $x: T_1, T_2$; $y: T_1, T_3$; $z: \emptyset$
- **C-edges:**
 - $T_1: R_x - T_2.W_x, W_x - T_2.W_x, R_y - T_3.W_y, W_y - T_3.W_y$
 - $T_2: R_x - T_1.W_x$ ($W_x - T_1.W_x$: see T_1)
 - $T_3: R_y - T_1.W_y$ ($W_y - T_1.W_y$: see T_1)
- **Remove S-edges:** $T_1: R_x - W_x, R_y - W_y$; $T_2: R_x - W_x$;
 $T_3: R_y - R_z, R_z - W_y$
- **Final chopping:**
 - $T_{11} : R_x, W_x$; $T_{12} : R_y, W_y$
 - $T_2 : R_x, W_x$
 - $T_3 : R_y, R_z, W_y$

Reordering Transactions

- Commutative operations:

- changing the order does not change the semantics of the program
- example: $R(y), R(z), W(y \leftarrow y + z)$ and $R(z), R(y), W(y \leftarrow y + z)$ do the same thing

- Transaction chopping:

- changing the order of commutative operations may lead to better chopping
- responsibility of the programmer to verify that operations are commutative!

- Example: consider $T_3 : R_y, R_z, W_y$ of the previous example

- assume T_3 computes $y + z$ and stores the sum in y
- then R_y and R_z are commutative and can be swapped
- $T'_3 : R_z, R_y, W_y$ can be chopped: $T'_{31} : R_z, T'_{32} : R_y, W_y$