

Index tuning

Viet-Trung Tran SolCT

Credit: Pei Li – Database management and performance tuning

Outline

- Query type
- Index type
- Index data structures



Query types: Why?

- Different indexes are good for different query types.
- We identify categories of queries with different index requirements.

Query types [cont'd]

- Point query: returns at most one record
 - SELECT nameFROM EmployeeWHERE ID = 8478
- Multipoint query: returns multiple records based on equality condition
 - SELECT nameFROM EmployeeWHERE department = 'IT'
- Range query on X returns records with values in interval of X
 - SELECT nameFROM EmployeeWHERE salary >= 155000



Query types

- Extremal query: returns records with max or min values on some attributes
 - SELECT name
 FROM Employee
 WHERE salary = MAX(SELECT salary FROM Employee)
- Ordering query: orders records by some attribute value
 - SELECT *
 FROM Employee
 ORDER BY salary
- Grouping query: partition records into groups; usually a function is applied on each partition
 - SELECT dept, AVG(salary)
 FROM Employee
 GROUP BY dept



Query types

- Join queries: link two or more tables
 - Equality join:
 - SELECT Employee.ssnum
 FROM Employee, Student
 WHERE Employee.ssnum = Student.ssnum
 - Join with non-equality condition:
 - SELECT e1.ssnum
 FROM Employee e1, Employee e2
 WHERE e1.manager = e2.ssnum
 AND e1.salary > e2.salary

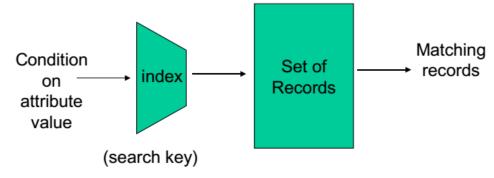


INDEX TYPE

What is an index?

An index is a data structure that supports efficient access

to data:



- Index tuning essential to performance! (2nd to Query tuning)
- Improper index selection can lead to:
 - indexes that are maintained but never used
 - files that are scanned in order to return a single record
 - multi-table joins that run for hours or days



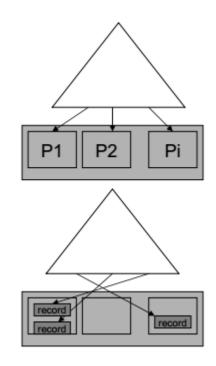
 Indexes can often be viewed as trees (B+-tree, hash) some nodes are in main memory (e.g., root) nodes deeper down in tree are less likely to be in main memory

Sparse index:

 pointers to disk pages at most one pointer per disk page usually much less pointers than records

Dense index:

pointers to individual records
 one key per record
 usually more keys than sparse index
 optimization: store repeating keys only once,
 followed by pointers



Spare vs. Dense

- Number of pointers:
 - ptrs in dense index = records per page × ptrs in sparse index
- Pro sparse: less pointers
 - typically record size is smaller than page size
 - less pointers result in less levels (and disk accesses)
 - uses less space
- Pro dense: index may "cover" query



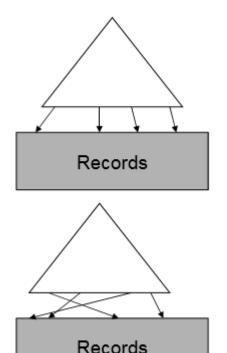
Covering index

- answers read query within index structure
 - fast, since data is not accessed
- Example:
 - dense index on lastname
 - SELECT COUNT(lastname) WHERE lastname='Smith'



Clustering vs. Non-Clustering

- Clustering index on attribute X
 - (also primary index)
 - records are grouped by attribute X on disk
 - B+-tree: records sorted by attribute X
 - only one clustering index per table
 - dense or sparse
- Non-clustering index on attribute X
 - (also secondary index)
 - no constraint on table organization more than one index per table always dense





Clustering index

Can be sparse:

fewer pointers than non-clustering index (always dense!)
 if record is small, save one disk access per record access

Good for multi-point queries:

- equality access on non-unique attribute
- all result records are on consecutive pages
- example: look up last name in phone book

Good for range, prefix, ordering queries:

- works if clustering index is implemented as B+-tree
- prefix example: look up all last names starting with 'St' in phone book
- result records are on consecutive pages



Clustering index [cont'd]

Good for equality join:

- fast also for join on non-key attributes
- index on one table: indexed nested-loop
- index on both tables: merge-join

Overflow pages reduce efficiency:

- if disk page is full, overflowing records go to overflow pages
- overflow pages require additional disk accesses



Equality Join with Clustering Index

- Example query:
 - SELECT Employee.ssnum, Student.course
 FROM Employee, Student
 WHERE Employee.firstname = Student.firstname
- Index on Emplyee.firstname:
 - use index nested loop join
 - for each student look up employees with same first name
 - all matching employees are on consecutive pages
- Index on both firstname attributes:
 - use merge join
 - read both tables in sorted order and merge (B+-tree)
 - each page read exactly once
 - works also for hash indexes with same hash function



Clustering Index and Overflow Pages

- Why overflow pages?
 - clustering index stores records on consecutive disk pages
 - insertion between two consecutive pages not possible
 - if disk page is full, overflowing records go to overflow pages
- Additional disk access for overflow page: reduced speed
- Overflow pages can result from:
 - Inserts
 - updates that change key value
 - updates that increase record size (e.g., replace NULL by string)
- Reorganize index:
 - invoke special tool
 - or simply drop and re-create index



Non-Clustering Index

- Always useful for point queries.
- Particularly good if index covers query.
- Multi-point query (but not covered): only if not too selective
 - nR: number of records returned by query
 - nP: number of disk pages in table
 - the nR records are uniformly distributed over all pages
 - thus query will read min(nR, nP) disk pages
- Index may slow down highly selective multi-point query:
 - scan is by factor 2–10 faster than accessing all pages with index
 - thus nR should be significantly smaller than nP



INDEX DATA STRUCTURES

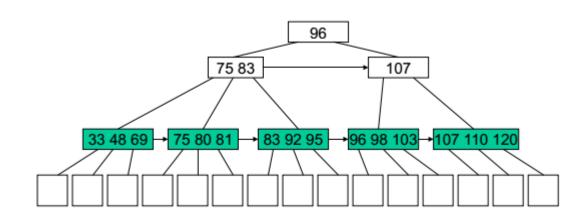
Index data structure

- Indexes can be implemented with different data structures.
 - B+-tree index
 - hash index
 - bitmap index (briefly)
 - dynamic hash indexes: number of buckets modified dynamically
 - R-tree: index for spacial data (points, lines, shapes)
 - quadtree: recursively partition a 2D plane into four quadrants
 - octree: quadtree version for three dimensional data
 - main memory indexes: T-tree, binary search tree



B+-Tree

- balanced tree of key-pointer pairs
- keys are sorted by value
- nodes are at least half full
- access records for key: traverse tree from root to leaf





Database Tuning

Key Length and Fanout

- Key length is relevant in B+-trees: short keys are good!
 - fanout is maximum number of key-pointer pairs that fit in node
 - long keys result in small fanout
 - small fanout results in more levels
 - Store 40M key-pointer pairs in leaf pages (page: 4kB, pointer: 4B)
 - 6B key: fanout 400 \Rightarrow 3 block reads per accesses

level	nodes	key-pointer pairs
1	1	400
2	400	160,000
3	160,000	64,000,000

• 96*B* key: fanout $40 \Rightarrow 5$ block reads per accesses

level	nodes	key-pointer pairs
1	1	40
2	40	1,600
3	1,600	64,000
4	64,000	2,560,000
5	2,560,000	102,400,000

• 6B key almost twice as fast as 96B key!



2/28/16

Estimate Number of Levels

- Page utilization:
 - examples assumes 100% utilization
 - typical utilization is 69% (if half-full nodes are merged)
- Number of levels:

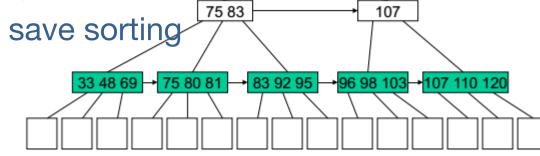
```
\begin{aligned} & \mathsf{fanout} = \lfloor \frac{\mathsf{node} \; \mathsf{size}}{\mathsf{key-pointer} \; \mathsf{size}} \rfloor \\ & \mathsf{number} \; \mathsf{of} \; \mathsf{levels} = \lceil \mathsf{log}_{\mathsf{fanout} \times \mathsf{utilization}}(\mathsf{leaf} \; \mathsf{key-pointer} \; \mathsf{pairs}) \rceil \end{aligned}
```

- Previous example with utilization = 69%:
 - -6B key: fanout = 400, levels = d3.11e = 4
 - -96B key: fanout = 40, levels = d5.28e = 6



B+-tree index supports

- point: traverse tree once to find page
- multi-point: traverse tree once to find page(s)
- range: traverse tree once to find one interval endpoint and follow pointers between index nodes
- prefix: traverse tree once to find prefix and follow pointers between index nodes
- extremal: traverse tree always to left/right (MIN/MAX)
- ordering: keys ordered by their value
- grouping: ordered keys save sorting





8/16 Databas

Hash index

Hash function:

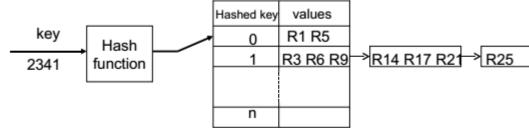
- maps keys to integers in range [0..n] (hash values)
- pseudo-randomizing: most keys are uniformly distributed over range
- similar keys usually have very different hash values!
- database chooses good hash function for you

Hash index:

- hash function is "root node" of index tree
- hash value is a bucket number

bucket either contains records for search key or pointer to overflow chain

with records



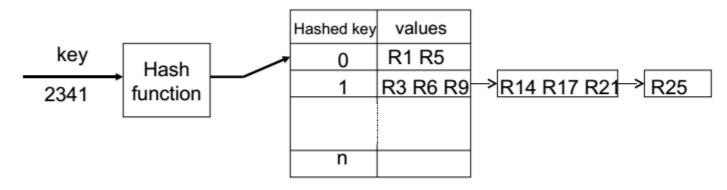
Key length:

- size of hash structure independent of key length
- key length slightly increases CPU time for hash function



Hash index usage

- Hash index supports
 - point: single disk access!
 - multi-point: single disk access to first record
 - grouping: grouped records have same hash value
- Hash index is useless for
 - range, prefix, extremal, ordering
 - similar key values have dissimilar hash values
 - thus similar keys are in different pages





Bitmap index

- Every value has it's own column == bitmap.
- One bit vector per attribute value (e.g., two for gender)
 - length of each bit vector is number of records
 - bit i for vector "male" is set if key value in row i is "male"

		$\pi_A(R)$
	1	3 2
Value List	2	2
value List	1 2 3 4 5 6 7 8	1
	4	2
	5	2 8 2 2 0 7
	6	2
	7	2
	8	0
	9	7
	10	5 6
	11	6
	12	4

B^8	B^7	B^6	B^5	B^4	B^3	B^2	B^1	${m B}^0$
0	0	0	0	0	1	0	0	0
0	0	0	0	0	0	1	0	0
0	0	0	0	0	0	0	1	0
0	0	0	0	0	0	1	0	0
1	0	0	0	0	0	0	0	0
0	0	0	0	0	0	1	0	0
0	0	0	0	0	0	1	0	0
0	0	0	0	0	0	0	0	1
0	1	0	0	0	0	0	0	0
0	0	0	1	0	0	0	0	0
0	0	1	0	0	0	0	0	0
0	0	0	0	1	0	0	0	0
				(b)				



Bitmap index

Works best if

- query predicates are on many attributes
- the individual predicates have high selectivity (e.g., male/ female)
- all predicates together have low selectivity (i.e., return few tuples)

Advantages

- Compact size.
- Efficient hardware support for bitmap operations (AND, OR, XOR, NOT).
- Fast search.
- Multiple differentiate bitmap indexes for different kind of queries.



Bitmap index example

 "Find females who have brown hair, blue eyes, wear glasses, are between 50 and 60, work in computer industry, and live in Bolzano"



INDEX TUNING EXERCISES

- The examples use the following tables:
 - Employee(ssnum, name, dept, manager, salary)
 - Student(ssnum, name, course, grade, stipend, evaluation)



Exercise 1 – Query for Student by Name

- Student was created with non-clustering index on name.
- Query:
 - SELECT *FROM StudentWHERE name='Bayer'
- Problem: Query does not use index on name.
- Solution: Try updating the catalog statistics.
 - Oracle, Postgres: ANALYZE
 - SQL Server: sp createstats
 - DB2: RUNSTATS



Exercise 2 – Query for Salary I

- Non-clustering index on salary.
- Catalog statistics are up-to-date.
- Query:
 SELECT *
 FROM Emplyee
 WHERE salary/12 = 4000
- Problem: Query is too slow.
- Solution: Index not used because of the arithmetic expression.
 - Rewrite query:
 - SELECT *
 FROM Emplyee
 WHERE salary = 48000



Exercise 3 – Query for Salary II

- Non-clustering index on salary.
- Catalog statistics are up-to-date.
- Query: SELECT * FROM Emplyee WHERE salary = 48000
- Problem: Query still does not use index. What could be the reason?
- How: The index is non-clustering. Many employees have a salary of 48000, thus the index may not help. It may still help for other, less frequent, salaries!

Exercise 4 – Clustering Index and Overflows

- Clustering index on Student.ssnum
- Page size: 2kB
- Record size in Student table: 1KB (evaluation is a long text)
- Problem: Overflow when new evaluations are added.
- Solution: Clustering index does not help much due to large record size. A non-clustering index avoids overflows.



Exercise 5 – Non-clustering Index I

- Employee table:
 - 30 employee records per page
 - each employee belongs to one of 50 departments (dept)
 - the departments are of similar size
- Query:
 - SELECT ssnumFROM EmplyeeWHERE dept = 'IT'
- Problem: Does a non-clustering index on Employee.dept help?
- Solution:
 - Only if the index covers the query.
 - 30/50=60% of the pages will have a record with dept = 'IT'
 - table scan is faster than accessing 3/5 of the pages in random order



Exercise 6 – Non-clustering Index II

- Employee table:
 - 30 employee records per page
 - each employee belongs to one of 5000 departments (dept)
 - the departments are of similar size
- Query:
 - SELECT ssnumFROM EmplyeeWHERE dept = 'IT'
- Problem: Does a non-clustering index on Employee.dept help?
- Solution:
 - Only if the index covers the query.
 - only 30/5000=0.06% of the pages will have a record with dept='IT'
 - table scan is slower



Exercise 7 – Statistical Analysis

- Auditors run a statistical analysis on a copy of Emplyee.
- Queries:
 - count employees with a certain salary (frequent)
 - find employees with maximum or minimum salary within a particular department (frequent)
 - find an employee by its social security number (rare)
- Problem: Which indexes to create?
- Solution:
 - non-clustering index on salary (covers the query)
 - clustering composite index on (dept, salary) using a B+-tree (all employees with the maximum salary are on consecutive pages)
 - non-clustering hash index on ssnum



Exercise 8 – Algebraic Expressions

- Student stipends are monthly, employee salaries are yearly.
- Query: Which employee is paid as much as which student?
- There are two options to write the query:
 - SELECT * SELECT * FROM Employee, Student WHERE salary = 12*stipend
 - SELECT * FROM Employee, Student WHERE salary/12 = stipend
- Index on a table with an algebraic expression not used.
- Problem: Which query is better?



Exercise 8 – Solution

- If index on only one table, it should be used.
- Index on both tables, clustering on larger table: use it.
- Index on both tables, non-clustering on larger table:
 - small table has (much) less records than large table has pages: use index on large table
 - otherwise: use index on small table



Exercise 9 – Purchasing Department

- Purchasing department maintains table
 - Onorder(supplier,part,quantity,price).
- The table is heavily used during the opening hours, but not over night.
- Queries:
 - Q1: add a record, all fields specified (very frequent)
 - Q2: delete a record, supplier and part specified (very frequent)
 - Q3: find total quantity of a given part on order (frequent)
 - Q4: find the total value on order to a given supplier (rare)
- Problem: Which indexes should be used?



Exercise 9 – Solution

Queries:

- Q1: add a record, all fields specified (very frequent)
- Q2: delete a record, supplier and part specified (very frequent)
- Q3: find total quantity of a given part on order (frequent)
- Q4: find the total value on order to a given supplier (rare)
- **Solution**: Clustering composite *B*+-tree index on (part, supplier).
 - eliminate overflows over night
 - attribute order important to support query Q3
 - hash index will not work for query Q3 (prefix match query)
- Discussion: Non-clustering index on supplier to answer query Q4?
 - index must be maintained and will hurt the performance of much more frequent queries Q1 and Q2
 - index does not help much if there are only few different suppliers



Exercise 10 – Point Query Too Slow

- Employee has a clustering B+-tree index on ssnum.
- Queries:
 - retrieve employee by social security number (ssnum)
 - update employee with a specific social security number
- Problem: Throughput is still not enough.
- Solution: Use hash index instead of B+-tree (faster for point queries).

