

# Transaction, ACID, concurrency tuning

Viet-Trung Tran

Information systems

Credit: Pei Li – Database management and performance tuning

# Outline

- Introduction to transaction
- ACID
- Concurrency control

# What is transaction

- A transaction is a unit of program execution that accesses and possibly updates various data items.
- Example: transfer \$50 from account  $A$  to account  $B$ 
  - 1.  $R(A)$
  - 2.  $A \leftarrow A - 50$
  - 3.  $W(A)$
  - 4.  $R(B)$
  - 5.  $B \leftarrow B + 50$
  - 6.  $W(B)$
- Two main issues:
  - 1. concurrent execution of multiple transactions
  - 2. failures of various kind (e.g., hardware failure, system crash)

# ACID Properties

- Database system must guarantee ACID for transactions:
  - **Atomicity**: either all operations of the transaction are executed or none
  - **Consistency**: execution of a transaction in isolation preserves the consistency of the database
  - **Isolation**: although multiple transactions may execute concurrently, each transaction must be unaware of the other concurrent transactions.
  - **Durability**: After a transaction completes successfully, changes to the database persist even in case of system failure.

# Atomicity

- Example: transfer \$50 from account  $A$  to account  $B$ 
  1.  $R(A)$
  2.  $A \leftarrow A - 50$
  3.  $W(A)$
  4.  $R(B)$
  5.  $B \leftarrow B + 50$
  6.  $W(B)$
- **What if failure (hardware or software) after step 3?**
  - money is lost
  - database is inconsistent
- **Atomicity:**
  - either all operations or none
  - updates of partially executed transactions not reflected in database

# Consistency

- Example: transfer \$50 from account  $A$  to account  $B$ 
  1.  $R(A)$
  2.  $A \leftarrow A - 50$
  3.  $W(A)$
  4.  $R(B)$
  5.  $B \leftarrow B + 50$
  6.  $W(B)$
- **Consistency in example:** sum  $A + B$  must be unchanged
- **Consistency in general:**
  - explicit integrity constraints (e.g., foreign key)
  - implicit integrity constraints (e.g., sum of all account balances of a bank branch must be equal to branch balance)
- **Transaction:**
  - must see consistent database
  - during transaction inconsistent state allowed
  - after completion database must be consistent again

# Isolation – Motivating Example

- Example: transfer \$50 from account  $A$  to account  $B$ 
  - 1.  $R(A)$
  - 2.  $A \leftarrow A - 50$
  - 3.  $W(A)$
  - 4.  $R(B)$
  - 5.  $B \leftarrow B + 50$
  - 6.  $W(B)$
- **Imagine second transaction  $T2$ :**
  - $T2 : R(A), R(B), \text{print}(A + B)$
  - $T2$  is executed between steps 3 and 4
  - $T2$  sees an inconsistent database and gives wrong result

# Isolation

- **Trivial isolation:** run transactions serially
- Isolation for concurrent transactions:
  - For every pair of transactions  $T_i$  and  $T_j$ ,
  - it appears to  $T_i$  as if either  $T_j$  finished execution before  $T_i$  started
  - or  $T_j$  started execution after  $T_i$  finished.
- **Schedule:**
  - specifies the chronological order of a sequence of instructions from various transactions
  - equivalent schedules result in identical databases if they start with identical databases
- **Serializable schedule:**
  - equivalent to some serial schedule
  - serializable schedule of  $T_1$  and  $T_2$  is either equivalent to  $T_1, T_2$  or  $T_2, T_1$



# Durability

- When a transaction is done it commits.
- **Example:** transaction commits too early
  - transaction writes *A*, then commits
  - *A* is written to the disk buffer
  - then system crashes
  - value of *A* is lost
- **Durability:** After a transaction has committed, the changes to the database persist even in case of system failure.
- **Commit only after all changes are permanent:**
  - either written to log file or directly to database
  - database must recover in case of a crash

# CONCURRENCY CONTROL

# Locks

- A lock is a mechanism to control concurrency on a data item.
- Two types of locks on a data item  $A$ :
  - **exclusive** –  $xL(A)$ : data item  $A$  can be both read and written
  - **shared** –  $sL(A)$ : data item  $A$  can only be read.
- Lock request are made to **concurrency control manager**.
- Transaction is blocked until lock is granted.
- **Unlock**  $A$  –  $uL(A)$ : release the lock on a data item  $A$

# Lock Compatibility

- Lock compatibility matrix:

$T_1 \downarrow T_2 \rightarrow$	shared	exclusive
shared	true	false
exclusive	false	false

- $T_1$  holds shared lock on  $A$ :**
  - shared lock is granted to  $T_2$
  - exclusive lock is not granted to  $T_2$
- $T_2$  holds exclusive lock on  $A$ :**
  - shared lock is not granted to  $T_2$
  - exclusive lock is not granted to  $T_2$
- Shared locks can be shared by any number of transactions.**

# Locking Protocol

- Example transaction  $T2$  with locking:
  1.  $sL(A), R(A), uL(A)$
  2.  $sL(B), R(B), uL(B)$
  3.  $print(A + B)$
- **$T2$  uses locking, but is not serializable**
  - $A$  and/or  $B$  could be updated between steps 1 and 2
  - printed sum may be wrong
- **Locking protocol:**
  - set of rules followed by all transactions while requesting/releasing locks
  - locking protocol restricts the set of possible schedules

# Pitfalls of Locking Protocols – Deadlock

- **Example:** two concurrent money transfers
  - $T1: R(A), A \leftarrow A + 10, R(B), B \leftarrow B - 10, W(A), W(B)$
  - $T2: R(B), B \leftarrow B + 50, R(A), A \leftarrow A - 50, W(A), W(B)$
  - possible concurrent scenario with locks:
  - $T1.xL(A), T1.R(A), T2.xL(B), T2.R(B), T2.xL(A), T1.xL(B), \dots$
  - $T1$  and  $T2$  block each other – no progress possible
- **Deadlock:** situation when transactions block each other
- **Handling deadlocks:**
  - one of the transactions must be rolled back (i.e., undone)
  - rolled back transaction releases locks

# Pitfalls of Locking Protocols – Starvation

- **Starvation:** transaction continues to wait for lock
- **Examples:**
  - the same transaction is repeatedly rolled back due to deadlocks
  - a transaction continues to wait for an exclusive lock on an item while a sequence of other transactions are granted shared locks
- Well-designed concurrency manager avoids starvation.

# Protocol: Two-Phase Locking

- Protocol that guarantees serializability.
- Phase 1: growing phase
  - transaction may obtain locks
  - transaction may not release locks
- Phase 2: shrinking phase
  - transaction may release locks
  - transaction may not obtain locks



# Two-Phase Locking – Example

- Example: two concurrent money transfers
  - $T1: R(A), A \leftarrow A + 10, R(B), B \leftarrow B - 10, W(A), W(B)$
  - $T2: R(A), A \leftarrow A - 50, R(B), B \leftarrow B + 50, W(A), W(B)$
- Possible two-phase locking schedule:
  - 1.  $T1 : xL(A), xL(B), R(A), R(B), W(A \leftarrow A + 10), uL(A)$
  - 2.  $T2 : xL(A), R(A), xL(B) \text{ (wait)}$
  - 3.  $T1 : W(B \leftarrow B - 10), uL(B)$
  - 4.  $T2 : R(B), W(A \leftarrow A - 50), W(B \leftarrow B + 50), uL(A), uL(B)$
- Equivalent serial schedule:  $T1, T2$

# CONCURRENCY TUNING

# Concurrency Tuning Goals

- **Performance goals:**
  - reduce blocking (one transaction waits for another to release its locks)
  - avoid deadlocks and rollbacks
- **Correctness goals:**
  - serializability: each transaction appears to execute in isolation
  - note: correctness of serial execution must be ensured by the programmer!
- Trade-off between performance and correctness!

# Ideal Transaction

- Acquires few locks.
- Favors shared locks over exclusive locks.
  - only exclusive locks create conflicts
- Acquires locks with fine granularity.
  - granularities: table, page, row
  - reduces the scope of each conflict
- Holds locks for a short time.
  - reduce waiting time

# Lock Tuning

1. Eliminate unnecessary locks
2. Control granularity of locking
3. Circumvent hot spots
4. Isolation guarantees and snapshot isolation
5. Split long transactions

# 1. Eliminate Unnecessary Locks

- Lock overhead:
  - memory: store lock control blocks
  - CPU: process lock requests
- Locks not necessary if
  - only one transaction runs at a time, e.g., while loading the database
  - all transactions are read-only, e.g., decision support queries on archival data

# 2. Control Granularity of Locking

- Locks can be defined at different granularities:
  - row-level locking (also: record-level locking)
  - page-level locking
  - table-level locking
- **Fine-grained locking (row-level):**
  - good for short online-transactions
  - each transaction accesses only a few records
- **Coarse-grained locking (table-level):**
  - avoid blocking long transactions
  - avoid deadlocks
  - reduced locking overhead

# Lock Escalation

- Lock escalation: (SQL Server and DB2 UDB)
  - automatically upgrades row-level locks into table locks if number of row-level locks reaches predefined threshold
  - lock escalation can lead to deadlock
- Oracle does not implement lock escalation.




# Granularity Tuning Parameters

- **1. Explicit control of the granularity:**
  - within transaction: statement within transaction explicitly requests a table-level lock, shared or exclusive (Oracle, DB2)
  - across transactions: lock granularity is defined for each table; all transactions accessing this table use the same granularity (SQL Server)
- **2. Escalation point setting:**
  - lock is escalated if number of row-level locks exceeds threshold (escalation point)
  - escalation point can be set by database administrator
  - rule of thumb: high enough to prevent escalation for short online transactions
- **3. Lock table size:**
  - maximum overall number of locks can be limited
  - if the lock table is full, system will be forced to escalate

# Overhead of table vs. row locking

- Experimental setting:
  - accounts(number, branchnum, balance)
  - clustered index on account number
  - 100,000 rows
  - SQL Server 7, DB2 v7.1 and Oracle 8i on Windows 2000
  - lock escalation switched off
- Queries: (no concurrent transactions!)
  - 100,000 updates (1 query)
    - example: update accounts set balance=balance\*1.05
  - 100,000 inserts (100,000 queries)
    - example: insert into accounts values(713,15,2296.12)



# 3. Circumvent Hot Spots

- Hot spot: items that are
  - accessed by many transactions
  - updated at least by some transactions
- Circumventing hot spots:
  - access hot spot as late as possible in transaction
  - (reduces waiting time for other transactions since locks are kept to the end of a transactions)
  - use partitioning, e.g., multiple free lists
  - use special database facilities, e.g., latch on counter

# Partitioning Example: Distributed Insertions

- **Insert contention:** last table page is bottleneck
  - appending data to heap file (e.g., log files)
  - insert records with sequential keys into table with B+ - tree
- **Solutions:**
  - use clustered hash index
  - if only B+ tree available: use hashed insertion time as key
  - use row locking instead of page locking

# Partitioning Example: DDL Statements and Catalog

- **Catalog:** information about tables, e.g., names, column widths
- Data definition language (DDL) statements must access catalog
- Catalog can become hot spot
- **Partition in time:** avoid DDL statements during heavy system activity

# 4. Weaken Isolation Guarantees

- The main idea is about consistency tuning
- Your homework

# 5. Chopping Long Transactions

- Algorithms for splitting long transactions
- Your homework